

CS 246E Lecture Notes
Instructor: Brad Lushman
Transcribed by: *Sibelius Peng*
FALL 2018

Contents

0.1. Readings	5
0.2. Motto	5
0.3. Sth... random?	5
I. Basic C++	6
1. Problem 1: Program Input and Output	8
1.1. Consider C	8
1.2. Consider C++	9
1.2.1. stdin/stdout	9
1.2.2. File access	10
1.3. References	12
1.4. Separate Compilation	13
1.5. Linux make	14
2. Problem 2: Linear Collection + Modularity	17
2.1. C preprocessor	18
2.1.1. Fixing the double-include problem	19
3. Problem 3: Linear Collection and Memory Management	21
3.1. Introduction to Classes	23
3.1.1. Initializing Objects	24
3.1.2. MIL	26
4. Problem 4: Copies	29
4.1. Copy and Swap Idiom	30
5. Problem 5: Moves	31
5.1. Copy/Move Elision	33
6. Problem 6: I want a constant vector	35
7. Problem 7: Tampering	37
8. Problem 8: Efficient Iteration	40
8.1. Iterator Pattern	40

9. Problem 9: Staying in Bounds	45
10. Problem 10: I want a vector of chars	49
11. Problem 11: Better initialization	51
12. Problem 12: I want a vector of Posns	54
13. Problem 13: Less Copying!	56
14. Problem 14: Memory Management is Hard!	60
15. Problem 15: Is vector exception safe?	64
16. Problem 16: Insert/remove in the middle	68
17. Problem 17: Abstraction Over Containers	70
18. Problem 18: Heterogeneous Data	73
19. Problem 19: I'm Leaking	79
20. Problem 20: I want a class with no objects	81
21. Problem 21: The copier is broken	82
22. Problem 22: I want to know what kind of Book I have	84
 II. A Big Unit on Objected-Oriented Design	 87
23. UML	88
24. Measures of Design Quality	90
24.1. Coupling & Cohesion	90
25. SOLID Principles of OO design	92
25.1. <u>S</u> ingle Responsible Principle	92
25.2. <u>O</u> pen-Closed Principle	93
25.3. <u>L</u> iskov Substitution Principle	94
25.4. <u>I</u> nterface Segregation Principle	98
25.5. <u>D</u> ependency Inversion Principle	102
26. Some More Design Patterns	105
26.1. Factory Method Pattern	105
26.2. Decorator Pattern	106
26.3. Visitor Pattern	111
 III. Abstraction in C++	 115
27. Problem 23: Shared OwnerShip	116
28. Problem 24: Abstraction over Iterators	119
28.1. Template Metaprogramming	123

29. Problem 25: I want an even faster vector	124
29.1. Move/Forward Implementation	127
30. Problem 26: Collecting Stats	130
31. Problem 27: Resolving Method Overrides At Compile-Time	132
32. Problem 28: Polymorphic Cloning	134
33. Problem 29: Logging	136
33.1. Mixin inheritance	137
34. Problem 30: Total Control	138
35. Problem 31: Total Control ove Vectors & Lists	141
36. Problem 32: A fixed-size allocator	142
37. Problem 34: I want a (tiny bit) smaller vector class	146
37.1. Something Missing from Last Year	146
IV. Other	148
38. Tutorials	149
38.1. Bash	149
38.2. Debugger	152
38.2.1. GDB Commands	152
38.3. Git	152
38.3.1. What is version control?	152
38.3.2. Using git: the basics	153
38.3.3. Working on multiple computers: setting up git	153
38.3.4. Preventing committing the wrong files	154
38.3.5. Merge conflicts	154
38.3.6. Retrieve old versions of files	154
38.3.7. Collaborating and rolling back more effectively: branches	155
38.3.8. ...and much more	155
38.4. Parsing	155
38.5. MVC	156
38.5.1. Ncurses	156
39. Readings	157

- DC 3110
- Note: tut is mandatory (I need to change my plan, tbh...)

0.1. Readings

§2.2, §4.3

0.2. Motto

C++ is easy because it is hard.

0.3. Sth... random?

- Linux - must use!
- Windows (don't care) - download putty.exe - connect ... (lazy to type) - enable X11 forwarding
- Mac/Linux - terminal (lazy to type, already familiar)
 - Also install an XWindows server, e.g. Xming(win), XQuartz(Mac) (I don't know why it doesn't work for me...)
- Xclock/Xeyes

Part I.

Basic C++

1. Problem 1: Program Input and Output	8
1.1. Consider C	8
1.2. Consider C++	9
1.2.1. stdin/stdout	9
1.2.2. File access	10
1.3. References	12
1.4. Separate Compilation	13
1.5. Linux make	14
2. Problem 2: Linear Collection + Modularity	17
2.1. C preprocessor	18
2.1.1. Fixing the double-include problem	19
3. Problem 3: Linear Collection and Memory Management	21
3.1. Introduction to Classes	23
3.1.1. Initializing Objects	24
3.1.2. MIL	26
4. Problem 4: Copies	29
4.1. Copy and Swap Idiom	30
5. Problem 5: Moves	31
5.1. Copy/Move Elision	33
6. Problem 6: I want a constant vector	35
7. Problem 7: Tampering	37
8. Problem 8: Efficient Iteration	40
8.1. Iterator Pattern	40
9. Problem 9: Staying in Bounds	45
10. Problem 10: I want a vector of chars	49

11. Problem 11: Better initialization	51
12. Problem 12: I want a vector of Posns	54
13. Problem 13: Less Copying!	56
14. Problem 14: Memory Management is Hard!	60
15. Problem 15: Is vector exception safe?	64
16. Problem 16: Insert/remove in the middle	68
17. Problem 17: Abstraction Over Containers	70
18. Problem 18: Heterogeneous Data	73
19. Problem 19: I'm Leaking	79
20. Problem 20: I want a class with no objects	81
21. Problem 21: The copier is broken	82
22. Problem 22: I want to know what kind of Book I have	84

Problem 1: Program Input and Output

Running a program from the command line:

```
./program-name # . means current directory      or  path/to/program-name
```

Providing inputs in two ways:

1. `./program-name arg1 arg2 ... argn`
Args are written into the program's memory (see Figure 1.1a)
2. `./program-name` Then type sth else
Input comes through standard input stream (stdin) instead (see Figure 1.1b)

Can be connected to other sources / destinations, e.g. files

Redirection `./my-prog < infile > outfile 2> errfile`

1.1. Consider C

```
#include <stdio.h>
void echo(FILE *f) {
    for (int c = fgetc(f); c != EOF; c = fgetc(f)) // fetch from file
        putchar(c); // put to stdout
}
int main(int argc, /* # command line args (always > 1) \count) */
        char *argv[] /* command-line args \vector)*/ ) {
    // argv[0] = program name ... argv[argc] = NULL
    if (argc == 1) echo(stdin);
    else
        for (int i = 1; i < argc; ++i) {
            FILE *f = fopen(argv[i], "r");
            echo(f);
            f(close);
        }
    return 0; // status code  convention -0 = OK
}
```

`echo` \$? status code of most recently run program.

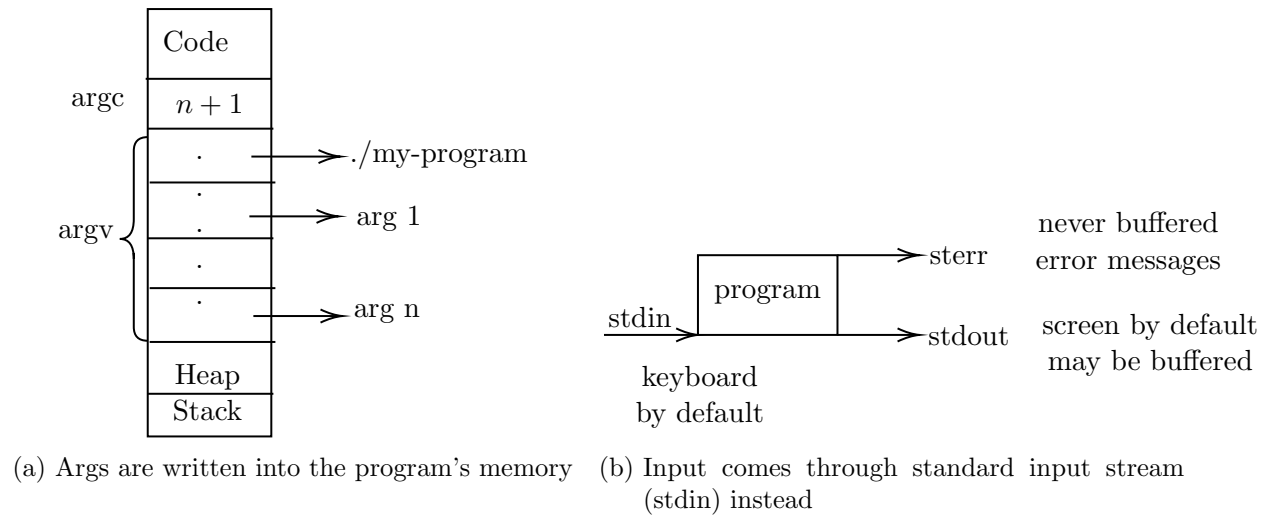


Figure 1.1.: two figures

Observe: command-line args/input from stdin - 2 different programming techniques.

To compile: `gcc -std=c99 -Wall myprogram.c -o myprogram` translates C source to executable binary

- `./myprogram`
- `Wall` (warn all) best friend
- By default `a.out`
- This program - simplification of the Linux `cat` command
 - `cat file1 file2 file3`
cat opens file1, file2, file3 prints contents one after the other
 - `cat`
- echoes stdin
 - `cat < file1`
 - * `file1` used as a source for stdin
 - * the shell (not cat) opens the file
 - * displays `file1`

`^D` (ctrl-d) signals end-of-file

1.2. Consider C++

Can we write the cat program in cpp? -already valid cpp

The "cpp way" : command line args - same as C

1.2.1. stdin/stdout

```
#include <iostream>
int main() {
    int x, y;
    std::cin >> x >> y;
```

```
std::cout << x+y << std::endl;
}
```

std::cin, std::cout, std::cerr - streams types.

- std::istream (std::cin)
 - std::ostream (std::cout, std::err)
-
- >> input operator cin >> x - populates x as a side-effect - returns cin
 - << output operator cout << x+y - prints x+y as a side-effect - returns cout

1.2.2. File access

```
#include <fstream>
std::ifstream f{"name-of-file"}; // initialization syntax - also opens the
    file
//ofstream for output
char c;
while (f >> c) { // evaluates to f
    // - implicitly converted to bool
    // - true if the read succeeded - false if failed
    std::cout << c;
}

// fclose is called when f goes out of scope - automatic
```

Recall: File access (identical to above)

```
Input: the quick brown fox
Output: thequickbrownfox
```

Stem input skips whitespace (just like scanf)

To include whitespace:

```
std::ifstream f{"name-of-file"};
f >> std::noskipws; // io manipulator
char c;
while (f >> c) {
    std::cout << c;
}
```

Note(important):

- no explicit calls to `fopen/fclose`
- initializing `f` opens the file
- when `f`'s scope ends, file is closed

Let's try to write a cat in C++

```
#include <iostream>
#include <fstream>
using namespace std; // Avoids having to say std::
void echo(istream f) {
    char c;
    f >> noskipws;
    while (f >> c) cout << c;
}

int main (int argc, char *argv[]) {
    if (argc == 1) echo(cin);
    else for (int i = 1; i < argc; ++i) {
        ifstream f {argv[i]};
        echo(f);
    }
}
```

Doesn't work - won't even compile - Why?

- `cin` - `istream` - `echo` takes an `istream`
- `f` - `ifstream` - is `echo(f)` a type mismatch?

No! - this is actually fine. `ifstream` is a *subtype* of `istream`

Any `ifstream` can be treated as an `istream`.

- foundational concept in OOP
- details later

That's not error.

The error is - you can't write a function that takes an `istream` the way `echo` does.

Why not? - time for discussion.

Compare:

```
int x;
scanf("%d", &x);

int x;
cin >> x;
```

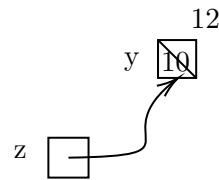


Figure 1.2.: z cannot change

C and C++ are pass-by-value languages - so `scanf` needs the address of `x`. in order to change `x`, via pointers.

So why not `cin >> &x`?

C++ has another ptr-like type.

1.3. References

See this code in Figure 1.2

```
int y = 10;
int &z = y; // z is an lvalue reference to y
// similar to
int * const z = &y; // but with auto-deferencing

z = 12; // NOT *z = 12;
// y now 12
int *p = &z; //gives the address of y
```

In all cases, `z` acts as if it were `y`.
`z` is an alias for `y`.

lvalue refs must be initialized, and must be initialized to something that has an address.

```
int &z = y; // OK
// int &x = 4 //wrong
// int &x = a+b; // wrong
```

- Cannot: - create a ptr a ref: `int &* x;` wrong
- Can create a ref to ptr: `int *&x = ...;` legal
- Can't create a ref to a ref: `int &&r = z;` wrong (will compile - means sth different)
- can't an array of refs: `int &r[3] = ...;` wrong

Can: pass as function parameters:

```
void inc(int &n) /* const ptr to the arg(x). Thus changes affect x */ {
    n = n + 1; // no ptr deref
}

int x = 5;
```

```
inc(x);

cout << x; // 6
```

`cin >> x` works because `x` is passed by ref.
`istream & operator>> (istream &in, int &x);`

Now consider `struct ReallyBig ...`

- `int f(ReallyBig rb) {...}`
 note - don't have say `struct ReallyBig` in C++
 pass-by-value - copies the struct (slow)
- `int g(ReallyBig &rb) {...}`
 reference - no copy (fast)
 but - could propagate changes to the caller
- `int h (const ReallyBig &rb) {...}`
 fast - no copy- `h` can't change `rb`

Prefer pass-by-const-ref over pass-by-value for anything larger than a ptr, unless the function needs to make a copy anyway - then use pass-by-val.

Also `int f(int &n) {...}`, `int g(const int &n) {...}`

- `f(5)`; X - can't initialize an lvalue ref (`n`) to a literal value (`5`)
- `g(5)`; OK - since `n` can never be changed, the compiler allows this. (stored in a temporary location, so `n` has sth to point at.)

Back to `cat`

- `f` passed by value
- `istream` is copied
- copying streams is not allowed
- Works if you pass `istream` by ref: `void echo(istream &f) { /*as before*/ }`

To compile: `g++ -std=c++14 -Wall myprogram.c -o myprogram`
`./myprogram`

1.4. Separate Compilation

Now we put `echo` in its module

```
// echo.h
void echo(istream &f);
//echo.cc
#include "echo.h"
void echo (istream &f) {...}
//main.cc
#include <iostream>
```

```
#include <fstream>
#include "echo.h"
int main(...) {
    ... echo(cin); ...
    ... ifstream f{argv[i]};
        echo(f);
    ...
}
```

Compile separately:

- `g++14 echo.cc` fails - no main
- `g++14 main.cc` fails - no echo

These give us linking errors.

Correct:

- `g++14 -c echo.cc` creates `echo.o`
- `g++14 -c main.cc` creates `main.o`

These are object files.

Linker: `g++14 echo.o main.o -o myprogram`

Advantage only have to compile the parts you change to relink

- change `echo.cc` - recompile `echo.cc`
- change `main.cc` - recompile `main.cc`
- `echo.h` - must recompile both (which includes this header)

And relink.

Change `echo.cc` – recompile `echo.cc`

Change `main.cc` – recompile `main.cc`

Change `echo.h` – recompile both `main.cc` and `echo.cc`

because they include `echo.h`, and link

What if we don't remember what changed. Or what depends on what?

1.5. Linux make

Create a file called Makefile

Contents:

```
my-program: main.o echo.o
    g++ main.o echo.o -o myprogram
main.o: main.cc echo.h
    g++ -std=c++14 -Wall -c main.cc
echo.o: echo.cc echo.h
    g++ -std=c++14 -Wall -c echo.cc
PHONY:clean
```

```
clean:
    rm myprogram *.o
```

targets: dependencies

\t recipes for building targets from dependencies (Note, can't use aliases)

From the command line: make (builds the whole program) Change just echo.cc, make (recompile echo.cc, relink)

make clean (remove all binaries)

How? List a directory in "Long form":

```
$ ls -l
-rw-r-----1  j2smith j2smith 25   Sep9 15:27 echo.cc
```

```
type |owner/group/other permissions| owner | group |size|last modified date/
time|name
```

based on last modified time.

Starting at the leaves of the dependency graph

If the dependency is newer than the target, rebuild the target ... and so on, up to the root target.

e.g echo.cc newer than echo.o – rebuild echo.o

echo.o newer than myprogram – rebuild myprogram

Shoutcuts use variables

```
CXX = g++                                # name of compiler
CXXFLAG = -std=c++14 -Wall               # compiler options
EXEC = myprogram
OBJECTS = main.o echo.o

${EXEC}: ${OBJECTS}
    ${CXX} ${OBJECTS} -o ${EXEC}
main.o: main.cc echo.h
echo.o: echo.cc echo.h                  # omit recipes -make "guess" the right one.
PHONY:clean
clean:
    rm ${OBJECTS} ${EXEC}
#Writing dependncies still hard, but compiler can help.

g++14 -c -MMD echo.cc
# generates echo.o echo.d
cat echo.d
echo.o: echo.cc echo.h

#Just include .d file into Makefile
```

```
CXX = g++                                # name of compiler
CXXFLAG = -std=c++14 -Wall -MMD # compiler options
EXEC = myprogram
OBJECTS = main.o echo.o
DEPENDS = ${OBJECTS:.o=.d}
${EXEC}: ${OBJECTS}
    ${CXX} ${OBJECTS} -o ${EXEC}
-include ${DEPENDS}
PHONY:clean
clean:
    rm ${OBJECTS} ${DEPENDS} ${EXEC}
```

Always use Makefile

Create the Makefile before you starting coding.

Problem 2: Linear Collection + Modularity

Linear Collection : linked lists + arrays

```
// node.h
struct Node {
    int data;
    Node *next;
};
size_t size(Node *n);

//node.cc
#include "node.h"

size_t size(Node *n){
    size_t count = 0;
    for(Node *cur = n; cur; cur = cur->next) ++count;
    return count;
}

// Note: do not use malloc / free in C++. Use new / delete
// main.cc
int main(){
    Node *n = new Node;
    n->data = 3;
    n->next = nullptr;
    Node *n2 = new Node {2, nullptr};
    Node *n3 = new Node {4, new Node{5, nullptr}};
    delete n;
    delete n2;
    // delete n3->next;
    // delete n3;
    while(n3){
        Node *tmp = n3;
        n3 = n3->next;
        delete tmp;
    }
}

// What happens if we do the following?
```

```
#include "node.h"
#include "node.h"
int main(){...}
```

Won't compile, struct is defined twice.
How can we prevent this?

2.1. C preprocessor

Transform the program before the compiler sees it.

`#include ...` drops the Contents of a file "right here"

Including old C headers - `#include <stdio.h>` become `#include<cstdio>`

```
#define VAR VALUE
```

- preprocessor variables
- all occurrences of VAR replaced with VALUE.

Eg

```
#define MAX 10
int x[MAX];
//OR
//----myprogram.cc-
int main(){
    int x[MAX];
    ...
}
```

`g++14 -D MAX=10 myprogram.called` //make a `#define` on the cmd line

```
// conditional compilation choose <= 1 of these to present to the compiler
#if SECURITYLEVEL == 1
short int
#elif SECURITYLEVEL == 2
long long int
#endif
public key;
```

special cases: -industrial-strength "comment out"

```
#if 0
...
#endif
```

However,

```

/*... */           // -----> doesn't nest
/*.. */           /* ... */
| only here are commented out |

```

2.1.1. Fixing the double-include problem

include guard

```

// node.h
# ifndef NODE_H // true if NODE_H is not defined
# define NODE_H // value is the empty string
... file
# endif

```

First time `node.h` is included - symbol `NODE_H` is not defined, so file is subsequently - symbol is defined, so file is suppressed.

Always put `#include` guards in header file.

Never compile `.h` files and include `.cc` files

Now - What if someone writes?

```

// tree.h
struct Node {
    int data;
    Node *left, *right;
};
size_t size(Node *n); //size of the tree

```

Can't use both this and linked list in same program!

Solution: namespaces

```

// list.h:
namespace List {
    struct Node {
        int data;
        Node *next;
    };
    size_t size(Node *n);
} // no ; here

// tree.h
namespace Tree {
    struct Node {
        int data;
        Node *left, right;
    };
    size_t size(Node *n);
}

```

```
} // no ; here
```

```
// list.cc
#include "list.h"

size_t List::size(Node *n) {...} // scope resolution operator

// tree.cc
#include "tree.h"
namespace Tree {
    size_t size(Node *n) {...}
}

// main.cc
#include "list.h"
#include "tree.h"
int main() {
    List::Node *ln = new List::Node {1, nullptr};
    Tree::Node *tn = new Tree::Node {1, nullptr};
    ..
    delete ln;
    delete tn;
}
```

Namespaces are open - anyone can add any namespace:

```
// some-other-file.h
namespace List {
    int some_other_f();
    struct some_other_struct {...};
}
```

Exception adding members to namespace std is not permitted. (you will succeed, compiler will allow it) (undefined behaviour, probably mess up)

Problem 3: Linear Collection and Memory Management

Read §7.7.1 §14 §16.2

```
// Arrays:
int a[10]; // on the stack, fixed size
// on the heap
int *p = new int[10];
// TO delete
delete [] p;
```

Use `new` with `delete`, use `new[]` with `delete[]`.
Mismatching these = undefined behaviour.

Problem what if our array isn't big enough?

Note no realloc for `new/delete`

Solution Use abstraction to solve the problem

```
#ifndef VECTOR_H
#define VECTOR_H
namespace CS246E {
    struct vector {
        int * thevector;
        size_t size, cap;
    };
    const size_t startsize = 1;
    vector makeVector();
    size_t size(const vector &v);
    int &itemAt(const vector &v, size_t i);
    // return a reference
    void push_back(vector &v, int n);
    void pop_back(vector &v);
    void dispose(vector &v);
}
#endif
//vector.cc
#include "vector.h"
```

```

namespace { //equivalent to static
    void increaseCap(CS246E::vector &v) {
        if (v.size == v.cap) {
            int newVec = new int[2*v.cap];
            for (size_t t=0; t<v.cap; ++t)
                newVec[t]=v.theVector[t];
            delete[] v.theVector;
            v.theVector = newVec;
            v.cap *= 2;
        }
    }
}

CS246E::vector CS246E::makeVector(){
    vector v { new int[startsize], 0, 1};
    return v;
}

size_t CS246E::size(const vector &v) {return v.size;}
int CS246E::itemAt(const Vector &v, size_t i){
    return v.theVector[i];
}

void CS246E::push_back(vector &v, int n) {
    increaseCap();
    v.theVector[v.size++]=n;
}

void CS246E::pop_back(vector &v){if(v.size>0) --v.size;}
void CS246E::dispose(vector &v) {delete[] v.theVector;}

// main.cc
#include "vector.h"
using CS246E::Vector;
int main() {
    Vector v = CS246E::makeVector(); // no arguments here
    push_back(v,1);
    push_back(v,10);
    push_back(v,100);
    // can add as many items as we want - never need to worry about
    allocation
    itemAt(v,0)=2;
    dispose(v);
}

```

Question why don't we have to say `CS246E::push_back`, `CS246E::itemAt`, `CS246E::dispose`?

Answer Argument-Dependent Lookup (ADL) - also called Koenig Lookup

If the type of the function `f`'s argument belongs to a namespace `n`, then C++ will search namespace `n`, as well as the current scope, for a function matching `f`.

Since `v`'s type (`Vector`) belongs to namespace `CS246E`, C++ will search namespace `CS246E` for `push_back`, `itemAt`, `dispose`

This is the same reason we can say `std::cout<<x`, rather than `std::operator<<(std::cout, x);`.

Problems

- what if we forget to call `makeVector`?
uninitialed object
- what if we forget to call `dispose`?
memory leak?

How can we make this more robust?

3.1. Introduction to Classes

First concept in OOP - functions inside of structs

```
struct Student {  
    int assns, mt, final;  
    float grade() { return assns*0.4+mt*0.2+final*0.4; }  
};
```

- structures that can contain functions - classes
- functions inside structs - called member functions or methods
- instances of a class - called objects

```
Student bob {90,70,80};  
// class object  
cout << bob.grade(); // method
```

What do `assns`, `mt`, `final` mean within `grade() {...}`

- fields of the current object - the receiver of the method call, e.g. `bob`

Formally - methods differ from functions in that methods take an implicit parameter called `this` that is a pointer to the receiver object.

`bob.grade();` → `this == &bob`

Could have written (equivalent):

```

struct Student {
    ...
    float grade() {
        return this->assns*0.4 + this->mt*0.2 + this->final*0.4;
    }
};

```

3.1.1. Initializing Objects

```

Student Bob {90, 70, 80};
// c-style struct initialization.
// field-by-field
// Ok, but limited.

```

Better - initialization method - a constructor

```

struct Student {
    int assns, mt, final;
    Student(int assns, int mt, int final) {
        this->assns = assns;
        this->mt = mt;
        this->final = final;
    }
};
Student bob {90, 70, 80};
// Now calls the ctor with args 90,70,80

```

Note once a ctor is defined, c-style field-by-field initialization is no longer available.

Equivalent Student bob = Student {90,70,80};

Heap Student *p = new Student {90,70,80};

Advantages of ctors

- automatically called
- default parameters
- overloading
- sanity checks

Eg

```

struct Student {
    ...
    Student (int assns=0, int mt=0, int final=0) { ... }
};

```



```
Student laura {70}; // 70,0,0
Student newkid; // 0,0,0
```

Note Every class comes with a built-in default (zero-arg) ctor

```
Node n; // calls default ctor
// - default constructs all fields that are objects
// - does nothing in this case
// goes away if you provide any ctor

// EG:
struct Node {
    int data;
    Node *next;
    Node (int data, Node *next = nullptr) {
        this->data = data; // etc
    }
};
Node n {3}; // 3, nullptr
// Node n; // wrong - no default ctor
```

Object creation protocol when an object is created
4 steps

1. space is allocated
2. later
3. Fields constructed in declaration order
(fields ctors are called for fields that are objects)
4. ctor body runs

Field initialization should happen in step3, but ctor body is in step4.

Consequence object fields could be initialized twice

```
// EG
#include <string>
struct Student {
    string name; //object
    Student(const string &name) {
        this->name = name;
    }
};
Student mike {"mike"};
// name default-initialized("") in step3 then reassigned in step4
```

To fix: Member Initialization List (MIL)

3.1.2. MIL

```
struct Student {
    ...
    Student(const string &name, int assns, int mt, int final):
name{name},assns{assns},// must be field names - before{}
    mt{mt},final{final} //step3
    // inside {} following usual scoping rules
    {} // Step4
};
```

MIL Must be used for fields that are

- constants
- references
- objects without a default ctor
- should be used as much as possible

Careful single arg ctors

```
struct Node {
    ...
    Node (int data, Node *next=nullptr): data{data}, next{next} {}
};
```

Single-arg ctors create implicit conversions

```
Node n{4}; // OK
Node n=4; // OK - int implicitly converted to Node
void f(Node n);
f(4); // OK - maybe trouble
f(x);
```

To fix

```
struct Node {
    ...
    explicit Node (...):...{} // disable the implicit conversion
};
Node n{4}; //OK
// Node n=4;
// f(4);
f(Node{4});
```

Object Destruction

- a method called the destructor (dtor) runs automatically
- built-in dtor - calls dtors on all fields that are objects

Object destruction protocol 4 steps

1. dtor body runs
2. fields destruction (dtors called on fields that are objects) in reverse declaration order
3. later
4. space deallocated

```
struct Node {
    int data;
    Node *next;
};
```

What does the built-in dtor do?

- nothing - neither fields is an object

```
Node *n = new Node {3, new Node{4,new Node {5,nullptr}}};
delete n; // only deletes the first node
//writing our own dtor:
struct Node {
    ...
    ~Node() { delete next; }
};
delete n; // now frees the whole list
```

Also

```
{
Node n{1,new Node {2, nullptr}};
// stack ----->|
//    |heap-----|
} // scope of n ends; n's dtor will run; whole list is freed
```

Objects

- a ctor always runs when they are created.
- a dtor always runs when they are destroyed

Back to vector problem - turn into a class:

```
//vector.h
#ifndef VECTOR_H
#define VECTOR_H
namespace CS246E {
    struct Vector {
        size_t n, cap;
        int *theVector;
        Vector();
        size_t size();
        int &itemAt(size_t i);
        void push_back(int n);
        void pop_back();
        ~Vector();
    };
}
#endif

//vector.cc
namespace {
    void increaseCap (const Vector &v) {...}
}

CS246E::Vector::Vector(): tsize{0}, cap{start_size}, theVector{new int[cap]}
{}
size_t CS246E::Vector::size() {return n;}
CS246E::Vector::~~Vector() {delete[] theVector;}

// main.cc
int main() {
    Vector v; // default ctor automatically called - no makeVector
    v.push_back(1);
    v.push_back(10);
    v.push_back(100);
    v.itemAt(0)=2;
    // no dispose - dtor cleans v up
}
```

Problem 4: Copies

```

Vector v;
v.push_back(100);
...
Vector w=v; // Allowed. Constructs w as a copy of v
w.itemAt(0); //100
w.itemAt(0) = 200;
v.itemAt(0);
// 200 Shallow copy - basis field-by-field copying
//                      - v and w share data
Vector w = v;
// constructs w as a copy of v
// involves the copy ctor

struct Vector {
    ...
    Vector (const Vector &other); // copy ctor
}; // compiler-supplied-copy ctor - copies all fields (shallow copy)

```

If you want a deep copy - write your own copy ctor

```

struct Node {
    int data;
    Node *next;
    ...
    Node (const Node &other): data{other.data},
        next{other.next? new Node{*other.next}: nullptr}
    {}
};

Vector v;
Vector w;
w = v; // copy, but not a constructor
        // copy assignment operator
        // compiler - supplied - copies each field (shallow)
        // compiler-supplied-copies each field (shallow)
        // - leaks w's old data

```

Deep copy assignment:

```
struct Node { // vector-exercise
    int data; Node *next;
    ...
    Node &operator=(const Node &other) {
        data = other.data;
        delete next;
        next = other.next? new Node{*other.next}: nullptr;
        return *this;
    }
}; // wrong - dangerous!
```

Consider: `Node n{...}; n=n;`

Destroys `n`'s data and then tries to copy it.

Must always ensure that `operator=` works in the cases of self-assignment.

```
Node &Node::operator=(const Node &other) {
    if (this == &other) return *this;
    // as before
}
```

Alternative: copy and swap idiom

4.1. Copy and Swap Idiom

```
#include <utility>
struct Node {
    ...
    void swap (Node &other){
        using std::swap;
        swap(data,other.data);
        swap(next,other.next);
    }
    Node &operator=(const Node &other) {
        Node tmp=other;
        swap(tmp);
        return *this;
    }
};
// drawback: always takes a copy
```

Problem 5: Moves

Consider

```
Node plusOne(Node n) {
    for (Node *p = &n; p; p=p->next) ++p->data;
    return n;
}
Node n {1, new Node{2,nullptr}};
Node m = plusOne(n); // copy ctor
                    // but what is "other" here? reference to what?
```

- temporary object → created to hold the result of `plusOne`.
- other is a reference to this temporary
 - copy ctor deep-copies the temporary into m

But

- the temporary is just going to be discarded as soon as the start `Node m = plusOne(n);` is done
- wasteful to deep copy the temp
- why not just steal the data instead?
 - save the cost of a copy
- Need to be able to tell whether other is a ref to a temporary object or a standalone object.

Rvalue references `Node &&` is a reference to a temporary object (rvalue) of type `Node`

```
// version of the ctor that takes a Node &&
struct Node {
    ...
    Node (Node &&other); // called a move ctor
};
Node::Node (Node &&other): data{other.data}, next{other.next}
{other.next=nullptr;}
```

Similarly `Node m; m = plusOne(n);` assignment from temporary

```
// More assignment operator
struct Node {
    Node &operator=(Node &&other) {
        // steal other's data
        // destroy my old data
        // easy - swap without copy
        using std::swap;
        swap(data, other.data);
        swap(next, other.next);
        return *this;
    }
};
```

Can combine copy/move assignment

```
struct Node {
    ...
    // pass-by-value invokes copy ctor if arg is an lvalue
    // invokes move ctor (if there is one) if arg is an rvalue
    Node &operator=(Node other){
        // unified assignment operator
        swap(other); //copies only if the arg is an lvalue
    }
};
```

Note copy and swap can be expensive - hand-coded operations may do less copying.

But now consider

```
struct Student {
    string name;
    Student (const string &name): name{name} {}
    // copies arg into the field (uses string's copy ctor)
};

// What if name points to an rvalue?
struct Student {
    string name;
    Student(std::string name): // copies if arg was an lvalue. Moves if arg
was an rvalue
    name{name/*another copy*/}{}
    // name may come from an rvalue, but name is an lvalue
};

struct Student {
    string name;
    Student(string name):
        name{std::move(name)} // now a move construction
    //std::move - forces a value to be treated as an rvalue
};
```



```

    {}
};
struct Student {
    Student(const Student &&other): // move ctor
        name {other.name}{}
    Student &operator=(Student other) {
        // unified assignment
        name = std::move(other.name);
        return *this;
    }
};

```

If you don't define move operations, copy operations will be used.

If you do define them, they will replace copy operations when the arg is a temporary (rvalue).

5.1. Copy/Move Elision

```

Vector makeAVector() { return Vector {};} // basic ctor
Vector v = makeAVector(); // move ctor? copy ctor?

```

Try with g++: -just the basic ctor - no copy/move In some circumstances, the compiler is allowed to skip calling copy/move ctors (but doesn't have to)

makeAVector() writes its result directly into the space occupied by V in the caller, rather than copy it over.

```

//Eg
Vector v/*copy ctor*/ = Vector/* basic ctor */{};
// Formally a basic construction + a copy construction

Student bob {90,70,80};
Student bob = Student{90,70,80};
//But the compiler is required to elide this copy construction, so basic
    ctor only

//EG
void doSomething(Vector v) {...} // pass-by-value
                                // ->copy/move ctor
doSomething(makeAVector());
// result of makeAVector written directly into the parameter
// no copy

```

This is allowed, even if dropping ctor calls would change the behaviour of the program. (e.g. if the ctor print sth)

If you need all ctors to run: g++11 -fno-elide-constructors can slow your program considerably

In summary Rule of 5 (Big 5)

If you need to customize any one of

1. Copy constructor
2. Copy assignment
3. Destructor
4. Move constructor
5. Move assignment

then you usually need to customize all 5.

From now on, we will assume that `Node` and `Vector` have the Big 5 defined.

Problem 6: I want a constant vector

Say we want to print a vector:

```
ostream &operator <<(ostream &out, const Vector &v){
    for (size_t i=0; i<v.size(); ++i){
        out << v.itemAt(i) << " ";
    }
    return out;
}
```

Won't compile! - can't call `size_t` and `itemAt` on a `const` object - What if these methods change fields?

Since they don't - declare them `const`!

```
struct Vector {
    size_t size() const;
    int &itemAt(size_t i) const;
    ...
};
size_t Vector::size() const {return n;}
int &Vector::itemAt(size_t i) const {return the Vector[i];}
```

Now the loop will work.

But

```
void f(const vector &v) {
    v.itemAt(0) = 4;    // Works!! v not very const...
}
```

- `v` is a `const` object - cannot change `n`, `cap`, `theVector` (ptr)
- You can changed items pointed to by `theVector`

Can we fix this?

```
struct Vector {
    ...
```

```

    const int &itemAt(size_t i) const {
        return theVector[i];
    }
};

```

Now `v.itemAt(0)=4` won't compile if `v` is `const`, but it also won't compile if `v` isn't `const`.
To fix: `const` overloading

```

struct Vector {
    ...
    const int &itemAt(size_t i) const;
    // Will be called if the object is const
    int &itemAt(size_t i);
    // Will be called if object is non-const
};

inline const int &Vector::itemAt(size_t i) const { return theVector[i]; }
inline int &Vector::itemAt(size_t) { return theVector[i]; }
// (inline) suggests replacing the function call with the function body-
// saves the cost of function call

```

More idomatic

```

struct Vector {
    ...
    const int &operator[] (size_t i) const { return theVector[i]; }
    int &operator[] (size_t i) { return theVector[i]; }
    // method body inside class - implicitly declares the method inline
};

```

Problem 7: Tampering

```
Vector v;
v.cap = 100; // sets cap without allocating memory
v.push_back(1);
v.push_back(2); // undefined behaviour - may crash
```

Interfering with ADTs:

1. Forgery
 - creating an object without calling a ctor function
 - not possible once we wrote ctors
2. Tampering
 - accessing the intervals without using a provided interface function

What's big deal? - Invariants → statement will always be true about an abstraction
ADTs provide and rely on invariants

- -stacks- provide the invariant that the last item pushed is the first item popped.
- -vectors- rely on an invariant that elements 0 ... cap-1 are valid to access

Can't rely on invariants if the user can interfere - makes the program hard to reason about.

Fix encapsulation - sealing objects into "black boxes"

```
namespace CS246E {
    struct Vector {
    private:
        size_t n, cap;
        int *theVector;
        // visible only within the Vector class
    public:
        Vector();
        size_t size() const;
        void push_back(int n);
        ...
        // visible to all
    };
}
```

```
};
// If no access specifier given: public
}
// vector.cc
#include "vector.h"
namespace {
    void increaseCap(Vector &v){...}
}
// doesn't work anymore! Doesn't have access to v's internals
```

Try again

```
// vector.h
namespace CS246E {
    struct Vector {
    private:
        size_t n, cap;
        int *theVector;
    public:
        Vector();
        ...
        private:
        void increaseCap(); // now a private method
    // don't need anonymous namespace anymore
    };
}
```

structs

- public default access
- private default access would be better
- class - exactly like struct, except private default access

```
class Vector {
    size_t n, cap;
    int *theVector;
public:
    Vector();
private:
    void increaseCap();
};
// vector.cc same
```

```
Node n {3, nullptr}; // stack-allocated
Node m {4, &n}; // dtor for m will try to delete &n. Undefined Behaviour
```

There was an invariant - `next` is either `nullptr`, or was allocated by `new`.

How can we enforce this?

Encapsulate `Node` inside inside a "wrapper" class

```
class list {
    struct Node {
        int data;
        Node *next;
        // ... methods
    };
    // private nested class
    // - list available outside class List
    Node *theList;
    size_t n;

public:
    list(): theList{nullptr}, n{0} {}
    ~list() { delete theList; }
    size_t size() const {return n;}
    void push_front (int x) {
        theList = new Node {x, theList};
        ++n;
    }
    void pop_front() {
        if (theList) {
            Node *tmp = theList;
            theList = theList->next;
            tmp->next = nullptr;
            delete tmp;
        }
    }
    const int &operator[] (size_t i) const {
        Node *cur = theList;
        for (size_t j = 0; j < i; cur = cur->next);
        return cur->data;
    }
    int &operator[] (size_t i) {...}
};
```

Client cannot manipulate the list directly - no access to next ptrs - invariant is maintained.

Problem 8: Efficient Iteration

```

Vector v;
v.push_back();
//...
for (size_t i = 0; i < v.size(); ++i) {
    cout << v[i] << " "; // O(1)
}
// Array access - efficient - O(n) traversal

list l;
l.push_front(..);
// ..
for (size_t i = 0; i < l.size(); ++i) {
    cout << l[i] << " "; // O(i)
}
// O(n^2) traversal

```

No direct access to "next" ptrs - how can we do efficient iteration?

Design patterns (more later)

- Well-known solutions to well-studied problems
- Adapted to suit-current needs

8.1. Iterator Pattern

For efficient iteration over a collection without exposing underlying structure.

Idea Create a class that "remembers" where you are in the list -abstraction of a ptr
Implementation: C

```
for (int *p = arr; p != a + size; ++p) printf("%d\n", *p);
```

```
class list {
    struct Node {...};
```



```

    Node *theList;
public:
    class iterator{
        Node *p;
    public:
        iterator(Node *p): p{p} {}
        bool operator!=(const iterator &other) const { return p != other.p; }
        iterator &operator++() {
            p=p->next;
            return *this;
        }
        int &operator*() { return p->data; }
    };
    ..
    iterator begin() { return iterator {theList}; }
    iterator end() { return nullptr; }
};

```

Now

```

list l;
...
for (list::iterator it = l.begin(); it != l.end(); ++it) {
    cout << *it << " ";
} // O(n)

```

Question Should `list::begin`, `list::end` be `const` methods?

Consider

```

ostream &operator<<(ostream &out, const list &l) {
    for (list::iterator it = l.begin(); it != end(); ++it) {
        out << *it << ' ';
    }
    return out;
} // won't compile if begin end are not const

```

If `begin/end` are `const`

```

ostream &operator<<(ostream &out, const list &l) {
    for (list::iterator it = l.begin(); it != end(); ++it) {
        out << *it << ' ';
        ++it; // increment items in the list
    }
    return out;
}

```

- will compile

- shouldn't - the list is supposed to be const
- compiles because `operator*` returns a non-const ref
- should compile if `l` is non-const

Conclusion

- iteration over `const` is different from iteration over non-const
- make a second iterator class

```
class list {
    ...
public:
    class iterator {
        // exactly as befor
        int &operator*() const;
    };
    class const_iterator {
        Node *p;
    public:
        const_iterator(Node *p): p{p} {}
        bool operator!=(const const_iterator &other) const;
        const_iterator &operator++();
        const int &operator*() const { return p->data; } // const at the front
    };

    iterator begin() { return iterator{theList}; }
    iterator end() { return iterator{nullptr}; }

    const_iterator begin() const { return const_iterator{theList}; }
    const_iterator end() const { return const_iterator{nullptr}; }
};

ostream &operation<<(ostream &out, const list &l) {
    for (list::const_iterator it = l.begin(); it != l.end(); ++it) {
        out << *it << '␣';
    }
    return out; // works
}
list::const_iterator it != l.begin; // mouthful
```

We got a shorter version!!!

```
ostream &operation<<(ostream &out, const list &l) {
    for (auto it = l.begin(); it != l.end(); ++it) {
        out << *it << '␣';
    }
    return out;
}
```

```
auto x = expr;
```

- saves writing x's type
- x is given the same type as expr

Emm, shorter...

```
ostream &operation<<(ostream &out, const list &l) {
    for (auto n : l) {
        out << n << ' ';
    }
    return out; // Range-based for loop
}
```

Range-based for loop requires

- methods (or functions) begin/end that return an iterator object
- the iterator class must support unary*, prefix++, !=

Note

```
for (auto n : l) ++n;
// n is declared by value, so ++n increments n, not the list items
for (auto &n : l) ++n;
// n is a ref - *will* update list items
for (const auto &n : l) ...
// n by ref (no copy) - cannot be modified
```

One small encapsulation problem...

Question What if client do `list::iterator it {nullptr};`

- forgery - client can create an end iterator without calling `l.end()`

To fix: make iterator ctors **private**

But then: list can't create iterators either

Solution: friendship

```
class list {
    ...
public:
    class iterator {
        ...
        iterator(Node *p);
    public:
        ...
        friend class list;
    };
    class const_iterator {
        ...
        const_iterator(Node *p);
    public:
```

```
    ...  
    friend class list;  
};  
    ...  
};
```

Limit friendships: - weaken encapsulation

Can do the same for vectors

```
class Vector {  
    ...  
public:  
    class iterator {...};  
    class const_iterator {...};  
    ...  
};
```

OR

```
class Vector {  
    size_t n, cap;  
    int *theVector;  
    ...  
public:  
    typedef int *iterator;  
    typedef const int *const_iterator;  
    iterator begin() { return iterator {theVector}; }  
    iterator end() { return iterator {theVector+n}; }  
    const_iterator begin() const { return const_iterator {theVector}; }  
    const_iterator end() const { return const_iterator {theVector+n}; }  
};
```

Problem 9: Staying in Bounds

```
Vector v;
v.push_back(2);
v.push_back(3);
v[2]; // Out of bounds! --- undefined behaviour - may be or may not crash
```

Can we make this safer?

Adding bounds checking to `operator[]` would be expensive.

Second, safer fetch operation:

```
class Vector {
...
public:
    int &at(size_t i) {
        if (i < n) return a[i];
        else // ... what???
    }
    ...
};
```

`v.at(2)`

- still not valid - what should happen?
- Returning any `int` - looks like a non-error
- Returning a non-`int` - not type-correct
- Crash the program - can we do better?

Don't return, don't crash

Solution raise an exception

```
class range_error {};
class Vector {
...
};
```

```
public:
    int &at(size_t i) {
        if (i < n) {...}
        else throw range_error {};
        // construct an object of type range_error and "throw" it
    }
};
```

Client options:

- 1) do nothing

```
Vector v;
v.push_back(1);
v.at(1); // the expression will crash the program
```

- 2) catch it

```
try {
    Vector v;
    v.push_back(1);
    v.at(1);
}
catch (range_error &r) { // catch by reference. - saves a copy
    operation           // r is the thrown object
    // do something
}
```

- 3) let your caller catch it

```
int f() {
    Vector v;
    v.push_back(1);
    v.at(1);
    return 0;
}
int g() {
    try {
        f();
    }
    catch (range_error &r) {
        // do sth
    }
}
```

Expression will propagate back through the call chain until a handler is found.

- called unwinding the stack
- if non handler found, program aborts
- control resumes after the catch block (problem code is not retried)

Question What happens if a ctor throws an expression?

- object is considered partially constructed.
- dtor will not run on partially constructed objects.

Eg

```
class C {...};
class D {
    C a;
    C b;
    int *c;
public:
    D () {
        c = new int[10];
        ...
        if (...) throw sth; // *
    }
    ~D () { delete[] c; }
};
D d;
```

At *, the D object is not fully constructed, so ~D() will not run on d.

But

- a & b are fully constructed, so their dtors will run
- So if a ctor wants to throw, it must clean up after itself

```
D () {
    c = new int[10];
    ...
    if (...) {
        delete[] c;
        throw something;
    }
}
```

Question What happens if a dtor throws? **Trouble**

- By default - program aborts immediately - `std::terminate` is called
- If you really want a throwing dtor - tag it with `noexcept(false)`
- but watch out

```
class myExn {};  
class C {  
    ...  
    ~C() noexcept(false) {  
        throw myExn {};;  
    }  
};  
void h() { C c1; } // dtor for c1 throws at the end of h  
void g() {  
    C c2;  
    h();  
    // - unwinding through g  
    // - dtor for c2 throws as we leave g  
}  
void f() { // no handler yet  
    try { g(); }  
    catch (myExn &e) /* handler */ {  
        ...  
    }  
}  
// now 2 unhandled expression  
// IMMEDIATE termination guaranteed (std::terminate)
```

Never let dtors throw. Swallow the expressions if necessary
Also note: you can throw any value - not just objects

Problem 10: I want a vector of chars

Start over? **NO!!!!**

Introduce a major abstraction mechanism - templates - generalize over types

```
template <typename T> class vector {
    size_t n, cap;
    T *theVector;
public:
    vector ();
    ...
    void push_back(T n);
    T &operator[] (size_t i);
    const T&operator[] const (size_t i);
    typedef T *iterator;
    typedef const T *const_iterator;
    ...
};

template <typename T> vector<T>::vector():
    n{0}, cap{1}, theVector{new T[cap]} {}

template <typename T> void vector<T>::push_back(T n) { ... }
// etc
```

Note must put the implementation in the .h file

```
// main.cc
int main() {
    vector <int> v; // vector of ints
    v.push_back(1);
    ...
    vector <char> w; // vector of chars
    w.push_back('a');
    ...
}
```

Semantics The first time the compiler encounters `vector<int>`, it creates a version of the vector class, where `int` replaces `T`, and then compiles the new class.

- can't do that unless it knows all the details about the class
- so implementation needs to be available, i.e. , in `.h`
- can also write method bodies in line.

11

Problem 11: Better initialization

Long sequence of `push_back` s is clunky.

```
// Arrays
int a[] = {1, 2, 3, 4, 5}; // :)
// Vector
vector<int> v;
v.push_back(1);
... // :(
```

Goal: better initialization

```
template<typename T> class Vector {
...
public:
    Vector() { ... }
    Vector(size_t n, T i = T {}):
        n{n}, cap{n==0 ? 1 : n}, theVector{ new T[cap] } {
        for (size_t j = 0; j < n; ++j) theVector[j] = i;
        }
    ...
};

// Now
Vector <int> v; // empty
Vector <int> {5}; // 0 0 0 0 0
Vector <int> v {3 , 5}; // 5 5 5
```

Note: `T {}` (default ctor) means 0 if `T` is a built-in type
Better - but what about true array-style initialization.

```
#include<initializer_list>
template <typename T> class Vector {
...
public:
    ...
    Vector() {...}
    Vector(size_t n, T i = T{}): ...
```

```

    Vector(std::initializer_list<T> init):
        n{init.size()}, cap{n==0 ? 1 : n}, theVector{ new T[cap] } {
            size_t i = 0;
            for (auto i: init) theVector[i++] = t;
        }
};

Vector<int> v {1 , 2, 3, 4, 5}; //1 2 3 4 5
Vector<int> v; // empty (default)
Vector<int> v {5}; // 5
Vector<int> v {3, 5}; // 3 5
// Oops - other constructor not being called

```

Default ctors take precedence over `initializer_list` ctors, which take precedence over all other ctors.

To get the other ctors to run - need round bracket initialization

```

Vector<int> v (5); // 0 0 0 0 0
Vector<int> v (3, 5); // 5 5 5

```

A note on cost: items in an `initializer_list` are stored in contiguous memory. (`begin` method returns a ptr)

So we are using one array to build another - 2 copies in memory

Note

- `initializer_lists` are meant to be immutable.
- do not try to modify their contents
- do not use them as standalone data structures

But also note, there is only one allocation in `Vector` - No doubling and reallocating, as there was with a sequence of `push_backs`

In general - if you know how big your `vector` will be, you can save reallocation cost by requesting the memory up front:

```

template<typename T> class Vector {
...
public:
    ...
    Vector() {...}
    ...
    void reserve(suize_t newCap) {
        if (cap < newCap) {
            T *newVec = new T[newCap];
            for (size_t i = 0; i < n; ++i) newVec[i] = theVector[i];
            delete[] theVector;
            theVector = newVec;
            cap = newCap;
        }
    }
}

```

```
    }  
};
```

Exercise rewrite `Vector` so that `push_back` uses `reserve` instead of `increaseCap`

```
Vector <int> v;  
v.reserve(10);  
v.push_back(...);  
...  
// can do 10 push_back without needing to reallocate
```

Problem 12: I want a vector of Posns

```
struct Posn {
    int x, y;
    Posn (int x, int y): x{x}, y{y} {}
};
int main() {
    Vector <Posn> v; // Won't compile! Why not?
}
```

Take a look at Vector's ctor:

```
template<typename T> vector<T>::Vector():n{0},cap{1},theVector{new
T {}} {}
creates an array of T objects
```

Which T object(s) will be stored in array?

C++ always calls a ctor when creating an object.

- Which ctor gets called? Default ctor T {}
- Posn doesn't have one

How can we create a vector of Posns?

Need to separate memory allocation (object creation step 1) from object initialization (step 2 - 4)

Allocation `void *operator new(size_t)`

- allocates `size_t` bytes - no initialization.
- returns a `void *`
- Note:
 - C: `void*` implicitly converts to any ptr type.
 - C++: the construction requires a cast

Initialization "placement new" new(addr) type

- constructs a "type" object at 'addr'
- does not allocate memory

```
template <typename T> class Vector {
...
public:
...
Vector(): n{0}, cap{1}, theVector{static_cast<T*>(operator new(sizeof(T))
)} {}
Vector(size_t n, T x): n{n}, cap{n==0? 1 : n},
    theVector{static_cast<T*>(operator new(cap*sizeof(T)))} {
    for (size_t i = 0; i < n; ++i) new (theVector+i) T(x);
}
...
void push_back(T x) {
    increaseCap();
    new (theVector +(n++)) T(x);
}
void pop_back() {
    if (n) {
        theVector[n-1].~T(); // Must explicitly invoke dtor.
        --n;
    }
}
~Vector() {
    destroyItems();
    operator delete (theVector);
}
void destroyItems() {
    for (T *p=theVector; p != theVector+n; ++p) p->~T();
    n=0;
}
};
```

13

Problem 13: Less Copying!

Before: `void push_back(int n);`

Now:

```
void push_back(T x)/*1*/ { // If T is an object, how many times is x being
    copied?
    increaseCap();
    new(theVector + n++) T(x); /* 2 */
}
```

If arg is an lvalue: ① is copy ctor ② is copy ctor 2 copies - want 1

If arg is an rvalue: ① is a move ctor ② is a copy ctor 1 copies - want 0

Fix

```
void push_back(T x) {
    increaseCap();
    new (theVector + n++) T (std::move(x));
}
```

lvalue: copy + move

rvalue: move + move

If T doesn't have a move ctor: 2 copies

Better: take T by reference:

```
void push_back(const T &x) { // no copy, no move
    increaseCap();
    new (theVector + n++) T(x); // copy ctor
}

void push_back(T &&x) { // no copy, no move
    increaseCap();
    new (theVector + n++) T(std::move(x)); // move ctor
}
```

lvalue: 1 copy rvalue 1 move

If T has no move ctor: 1 copy

Now consider

```
vector<Posn> v;
v.push_back(Posn{3,4});
```

- (1) ctor call to create the object
- (2) copy or move constructor into the vector (depending on whether Posn has a move ctor)
- (3) dtor call on the temp object

Could eliminate (1) and (3) if we could get vector to create the object, instead of the client

- pass ctor args to the vector, not the actual object

A note on template functions:

Consider: `std::swap` - seems to work on all types

Implementation:

```
template<typename T> void swap(T &a, T&b){
    T tmp(std::move(a));
    a = std::move(b);
    b = std::move(tmp);
}

int x = 1, y = 2;
swap(x, y); // Equiv: swap<int>(x, y);
```

Don't have to say `swap<int>`. C++ can deduce this from the types of `x` and `y`

In general, only have to say `f<T>(...)` if `T` cannot be deduced from args.

Type deduction for template args follows the same rules as type deduction for `auto`

Back to vectors - passing ctor args

- we don't know what types ctor args should have
- `T` could be any class, could have several ctors

Idea member template function - like `swap`, it could take anything

Second problem how many ctor args?

Solution variadic templates

```
template<typename T> class Vector {
//...
public:
    //...
    template<typename... Args>
    void emplace_back (Args... args) {
        increaseCap();
        new (theVector + n++) T(args...);
    }
};
```

```
    }
};
```

`Args` is a sequence of type variables denoting the types of the actual arguments of `emplace_back`

`args` is a sequence of program variables denoting the actual args of `emplace_back`

```
vector<Posn> v;
v.emplace_back(3, 4);
```

Problem `args` is taken by value
 Can we take `args` by reference?
 lvalue or rvalue reference?

```
template<typename... Args> void emplace_back(Arg &&... args) {
    // ...
}
```

Special rules here: `Args &&` is a universal reference

- can point to an lvalue or an rvalue

When is a reference universal? Must have the form `T &&`, where `T` is the type arg being deduced for the current template function call.

Examples

```
template<typename T> class C {
public:
    template<typename U> void f (u &&x); // universal
    template<typename U> void g (const U &&x); // not universal, just rvalue
    void h (T &&x); // not universal, T is not being deduced
```

Now recall

```
class C {...};
void f(C &&x) { // rvalue ref x points to an rvalue, but x is an lvalue
    g(x); // x is passed as an lvalue to g
}
```

If you want to preserve the fact that `x` is an rvalue ref, so that a "move" version of `g` is called (assuming there is one):

```
void f(C &&x) {
    g(std::move(x));
}
```

In the cases of args - don't know if args are lvalues, rvalues, or a mix. Want to call `move` on args iff the args are rvalues

```
template <typename... Args> void emplace_back(Args &&... args) {  
    increaseCap();  
    new (theVector + n++) T (std::forward<Args>(args) ...);  
}
```

`std::forward`

- calls `std::move` if its argument is an rvalue ref
- else does nothing

Now `args` is passed to `T`'s ctor with lvalue/rvalue information preserved - called perfect forwarding.

Problem 14: Memory Management is Hard!

No it isn't.

Vectors

- can do everything arrays can
- grow as needed - $O(1)$ amortized time or better
- clean up automatically when they go out of scope
- are tuned to minimize copying

Just use vectors, and you'll never have to manage arrays again.

C++ has enough abstraction facilities to make programming easier than C

But what about single objects?

```
void f() {  
    Posn *p = new Posn {1, 2};  
    ...  
    delete p; // must deallocate the Posn!
```

First ask: Do you really need to use the heap? Could you have used the stack instead?

```
void f() {  
    Posn p {1, 2};  
    ...  
} // no cleanup necessary
```

Sometimes you do need the heap. Calling `delete` isn't so bad. But consider:

```
class BadNews {};  
void f() {  
    Posn *p = new Posn {1, 2};  
    if (some condition) { throw BadNews{}; }  
    delete p;  
}
```

`p` is leaked if `f` throws. Unacceptable.

Raising and handling an exception should not corrupt the program - we desire exception safety

Leaks are a corruption of the program's heap - will eventually degrade performance and crash.

If a program cannot recover from an exception without corrupting its memory, what's the point of recovering?

What's constitutes exception safety? 3 levels:

- 1) Basis guarantee - once an exception has been handled, the program is in some valid state. No leaked memory, no corrupted data structure, all invariants maintained.
- 2) Strong guarantee - if an exception propagates outside a function `f`, then the state of the program will be as if `f` had not been called. `f` either succeeds completely or not at all.
- 3) Nothrow guarantee - a function `f` offers the nothrow guarantee if it never emits an exception and always accomplished its purpose.

Will revisit.

Coming back to `f`...

```
void f() {
    Posn *p = new Posn {1, 2};
    if (some condition) {
        delete p;
        throw BadNews {};
    }
    delete p;
}
// duplicated effort - memory management is even harder!
```

Want to guarantee that `delete p` happens no matter what.

What guarantees does C++ offer? Only one: that dtors for stack-allocated objects will be called when the objects go out of scope.

So create a class with a dtor `deletes` the ptr.

```
template<typename T> class unique_ptr {
    T *p;
public:
    unique_ptr(T *p): p{p} {}
    ~unique_ptr() { delete p; }
    T *get() const { return p; }
    T *release() {
        T *q = p;
        p = nullptr;
        return q;
    }
};

void f() {
    unique_ptr<Posn> p { new Posn {1, 2} };
}
```

```
if (some condition) { throw BadNews {}; }
}
```

That's it - less memory management effort than we started with.

Using `unique_ptr` - can use `get` to fetch the ptr.
Better - make `unique_ptr` act like a ptr.

```
template<typename T> class unique_ptr {
    T *p;
public:
    ...
    T &operator*() const { return *p; }
    T *operator->() const { return p; }
    explicit operator bool() const { return p; } // prohibits bool b = p;
    void reset(T *p1) {
        delete p;
        p = p1;
    }
    void swap (unique_ptr<T> &x) {
        std::swap(p, x.p);
    }
};

void f() {
    unique_ptr <Posn> p { new Posn {1, 2} };
    cout << p->x << " " << p->y;
}
```

But consider

```
unique_ptr <Posn> p { new Posn {1, 2} };
unique_ptr <Posn> q = p;
```

Two ptrs to the same object! Can't both delete it! Undefined behaviour.

Solution copying `unique_ptr`s is not allowed. Moving OK, though.

```
template <typename T> class unique_ptr {
    ...
public:
    ...
    unique_ptr (const unique_ptr &other) = delete;
    unique_ptr &operator (const unique_ptr &other) = delete;
    // this is how copying of streams is prevented
    unique_ptr (unique_ptr &&other): p{other.p} { other.p = nullptr; }
    unique_ptr &operator= (unique_ptr &&other) {
        swap(other);
        return *this;
    }
};
```

```
    }  
};
```

Small exception safety issue. Consider

```
class C {...};  
void f(unique_ptr<C> x, int y) {...}  
int g() {...}  
f(unique_ptr<C> {new C}, g());
```

C++ does not enforce order of argument evaluation. It could be

- (1) `new C`
- (2) `g()`
- (3) `unique_ptr<C> {(1)}`

- then what if `g` throws? (1) is leaked
- can fix this by making (1) and (3) inseparable. - use a helper function

```
template <typename T, typename... Args> unique_ptr<T> make_unique(Args&&...  
    args) {  
    return unique_ptr<T> { new T (std::forward<Args>(args)...) };  
}
```

Example becomes `f(make_unique<C>(), g());` - No leak if `g` throws. `unique_ptr` is an example of the C++ idiom Resource Acquisition Is Initialization (RAII)

- any resource that must be properly released (memory, file handle, etc) should be wrapped in a stack-allocated object whose dtor fixes it.

eg `unique_ptr`, `ifstream`

- acquire the resource when the object is initialized
- release it when the object's dtor runs

Problem 15: Is vector exception safe?

Consider:

```
template <typename T> class vector {
    size_t n, cap;
    T *theVector;
public:
    vector(size_t n, const T &x): n{n}, cap{cap},
        theVector{static_cast<T*>(operator new(n*sizeof(T)))} {
        for (size_t i = 0; i < n; ++i) new (theVector+i) T(x); // copy ctor -
        what if this throws?
    }
};
```

- partially constructed vector - dtor will not run
- broken invariant - the vector does not contain n valid objects.

Note if `operator new` throws - nothing has been allocated - no problem - strong guarantee
Fix:

```
template <typename T> vector <T>::vector(size_t n, const T &x): n{n}, cap{cap},
    theVector{/* as before */} {
    size_t progress = 0;
    try {
        for (size_t i = 0; i < n; ++i) {
            new (theVector + i) T(x);
            ++progress;
        }
    }
    catch (...) { // catch anything
        for (size_t i = 0; i < progress; ++i) theVector[i].~T();
        operator delete(theVector);
        throw; // rethrow
    }
}
```


Abstract the filling part into its own function:

```
template <typename T> void uninitialized_fill (T *start, T *finish, const T
    &x) {
    T *p;
    try {
        for (p = start; p != finish; ++p) new (p) T (x);
    }
    catch (...) {
        for (T *q = start; q != p; ++q) q->~T();
        throw;
    }
}

template <typename T> vector<T>::vector(size_t n, const T &x): n{n}, cap {n
    },
    theVector {static_cast<T*>(operator new(n*sizeof(T)))} {
    try {
        uninitialized_fill(theVector, theVector+n, x);
    }
    catch (...) {
        operator delete (theVector);
        throw;
    }
}
```

Can clean this up by using RAII on the array

```
template <typename T> struct vector_base {
    size_t n, cap;
    T *v;
    vector_base (size_t n): n{n}, cap{cap}, v{static_cast <T*> (operator new (n
        *sizeof(T)))} {}
    ~vector_base() { operator delete(v); }
};

template <typename T> class vector {
    vector_base <T> vb; // cleaned up implicitly when vector is destroyed
public:
    vector (size_t n, const T &x): vb{n} {
        uninitialized_fill(vb.v, vb.v+vb.n, x);
    }
    ~vector() { destroy_elements(); }
    void destroy_elements() {
        for (T *p = vb.v, p != vb.v+vb.n; ++p) p->~T();
    }
};
```

Strong guarantee

Copy ctor:

```
template <typename T> vector<T>::vector(const vector &other): vb{other.size
    ()} {
    uninitialized_copy (other.begin(), other.end(), vb.v);
    // similar to uninitialized_fill. Details - exercises
}
```

Assignment: copy and swap is exception safe - as long as swap is nothrow
push_back:

```
template <typename T> vector<T>::push_back (const T &x) {
    increaseCap(); // details - exercise ... done below
    /* new(vb.v + vb.n++) T(x); */
    // don't increment n before you know the construction has succeeded
    // instead
    new(vb.v + vb.n) T(x);
    ++vb.n;
    // if this throws - have the same vector
}
```

What if increaseCap throws?

```
void increaseCap() {
    if (vb.n == vb.cap) {
        vector_base vb2 {2 * vb.cap}; // RAI
        uninitialized_copy(vb.v, vb.v + vb.n, vb2.v); // Strong guarantee
        // change this into uninitialized_copy_or_move
        destroy_elements();
        std::swap(vb, vb2); // no throw on ints and ptrs
    }
}
```

Note: only try blocks we use are in uninitialized_fill and uninitialized_copy

Efficiency

- copying from the old array to the new one - moving would be better
- but moving destroys the old array so if an exception is thrown during moving, our vector is corrupted.

So we can only move if we are sure that move operation is nothrow.

```
template <typename T> void uninitialized_copy_or_move(T *start, T *finish, T
    *target) {
    T *p;
    try {
        for (p = start; p != finish; ++p, ++target) {
            new(p) T(std::move_if_noexcept(*target));
        }
    }
    catch (...) { // will never happen if T has a non-throwing move
    }
```

```

    for (T *q = start; q != p; ++q) q->~T();
    throw;
}
}

```

```
std::move_if_noexcept(x);
```

- produces `std::move(x)` if x has non-throwing move ctor
- produces x otherwise

But how should the compiler know if T's move ctor is non-throwing?
You have to tell it.

```

class C {
public:
    C (C &&other) noexcept;
};

```

In general, `move` and `swap` should be no-throwing.
Declare them so - will allow more optimization codes to run. Any functions you are sure will never have to throw or propagate an exception, you should declare `noexcept`.

Q: Is `std::swap` `noexcept`?

```

template <typename T> void swap (T &a, T &b) _____ {
    T c (std::move(a));
    a = std::move(b);
    b = std::move(c);
}

```

A: Only if T has a `noexcept` move ctor and move assignment. Thus blank should be filled with

```

noexcept(std::is_nothrow_move_constructible<T>::value &&
        std::is_nothrow_move_assignable<T>::value)

```

Note `noexcept` means the same as `noexcept(true)`

Problem 16: Insert/remove in the middle

A method like `vector<T>::insert(size_t y, const T &x)` is easy to write. The same method for `list<T>` would be slow (linear traversal). Using iterators can be faster.

```
template <typename T> class Vector {
public:
    iterator insert (iterator posn, const T &x) {
        increaseCap();
        // - add new items to end
        // - shuffle items down (move_if_noexcept)
        return iterator to point of insertion
    }
};
```

adding the codes...

```
template <typename T> class Vector {
public:
    iterator insert (iterator posn, const T &x) {
        ptrdiff_t offset = posn - begin();
        increaseCap();
        iterator newPosn = begin() + offset;
        new (end()) T (std::move_if_noexcept (* (end()-1)));
        ++vb.n;
        for (iterator it = end() - 1; it != newPosn; --it) {
            *it = std::move_if_noexcept (* (it-1));
        }
        *newPosn = x;
        return newPosn;
    }
};
```

Exception safe?

- assuming T's copy/move options are exception safe (at least basis guarantee)
- may get a partially shuffled vector, but it will still be a valid vector

Note if you have other iterators pointing at the vector:

1	2	3	4	...
it1		here	it2	

and you insert at 'here':

1	2	5	3	4	...
it1		here	it2		

it 2 points at different item.

Convention: after a call to insert or erase, all iterators pointing after the point of insertion/erasure are considered invalid and should not be used.

Similarly if a reallocation happens - all iterators pointing at the vector become invalid.

Exercises: erase

- remove the item pointed to by an iterator
- return an iterator to the point of erasure

emplace - like insert, but takes ctor args.

But - that means we have a problem with `push_back`. If placement new throws, the vector is the same, but iterators were invalidated!

To fix:

- allocate new array
- place the new item
- copy or move old items
- delete

`void pop_back()` Not `T pop_back()`.

- would call copy (or move) ctor on return
- then dtor on item in vector
- if copy ctor throws, dtor call will not happen
- not exception safe

above needs to be changed a little. If the copy ctor throws, the client doesn't get the item and the item is popped, and it can not recover gracefully.

Problem 17: Abstraction Over Containers

Recall: map from Racket

```
(map f (list a b c)) --> (list (f a) (f b) (f c))
```

May want to do the same with vectors

source

a	b	c	d
---	---	---	---

 \rightarrow target

f(a)	f(b)	f(c)	f(d)
------	------	------	------

```
template <typename T1, typename T2>
void transform(const vector <T1> &source, vector<T2> &target, T2 (*f)(T1)) {
    auto it = target.begin();
    for (auto &x : source) {
        *it = f(x);
        ++it;
    }
}
```

OK, but ..

- what if we only want part of source?
- what if we want to send the source to the middle of target, not the beginning?

More general: use iterators.

```
template <typename T1, typename T2>
void transform(vector <T1>::iterator start, vector <T1>::iterator finish,
              vector<T2>::iterator target, T2 (*f)(T1)) {
    while (start != finish) {
        *target = f(*start);
        ++target;
        ++start;
    }
} // ... won't compile
```

(Assume it works) OK, but

- if I want to transform a list, I will write the same code again.

- what if I want to transform a list a vector or vice versa?
- make the type variables stand for the iterators, not the container elements.
- but then how do we indicate the type of f?

```
template <typename InIter, typename OutIter, typename Fn>
void transform (InIter start, InIter finish, OutIter target, Fn f) {
    while (start != finish) {
        *target = f(*start);
        ++target;
        ++start;
    }
}
```

works over vector iterators, list iterators, or any other kinds of iterators.

InIter/OutIter can be any types that support ++, !=, *, including ordinary pointers.

C++ will instantiate this template function with any type that has the operations used by the function.

Function can be any type that supports function application.

Eg

```
class Plus {
    int n;
public:
    Plus (int n): n{n} {}
    int operator() (int m) { return m + n; }
};

Plus p {5};
cout << p(7); // 12
// function object

transform(v.begin(),v.end(),w.begin(), Plus{1}); // in last lecture of cs246
...
// OR
transform(v.begin(),v.end(),w.begin(), [](int n){return n+1;});
//                                     |<-   lambda   ->|
```

lambda:

```
[.../* capture list */...](.../* parameter list */...) mutable? noexcept? { .../*
    body */ ...}
```

semantics:

```
void f(T1 a, T2 b) {
    [a, &b] (int x) { return body; } (args)
```

```
}  
// translated to  
void f(T1 a, T2 b) {  
    class ~~~ { // name not known to you  
        T1 a;  
        T2 &b;  
    public:  
        ~~~ (T1 a, T2 &b): a{a}, b{b} {}  
        auto operator() (int x) const { return body; }  
    };  
    ~~~ {a, b}.operator() (args);  
}
```

If the lambda is declared mutable, then `operator()` is not `const`.

capture list - provides access to selected vars in the enclosing scope.

18

Problem 18: Heterogeneous Data

I want a mixture of types in my vector - can't do this with a template
eg

```
vector<template <typename T> T> v; // X
```

- not allowed - templates are compile-time entities - don't exist at run time.

eg fields of a struct

```
class MediaPlayer {  
    ...  
    template <typename T> T nowPlaying; // songs, movies etc  
    // can't do this
```

What's available in C:

```
// unions:  
union Media {Song s; Movie m;};  
Media nowPlaying;  
  
// void *:  
void *nowPlaying;  
  
// both not type-safe
```

Items in a heterogeneous collection will usually have something in common, e.g. provide a common interface.

Can be views as different "kinds" of a more general "thing".

So have a vector of "things" of a field of type "thing".

We'll use the standard CS 246 example (will be covered tomorrow in CS 246) :

```
class Book { // superclass, or base class  
    string title, author;  
    int length;  
public:  
    Book(string title, string author, int length): ... {}
```

```

    bool isHeavy() const { return length > 100; }
    string getTitle() const { return title; }
    // etc
};

```

Book	title
	author
	length

```

class Text: public Book { // subclass, derived class
    string topic;
public:
    Text(string title, string author, int length, string topic):
        Book{title, author, length}, topic{topic} {}
    // if we don't do this with book, default ctor will be invoked, but we
    // don't have one
    bool isHeavy const { return length > 200; }
    string getTopic const { return topic; }
};

class Comic: public Book {
    string hero;
public:
    Comic (...) ... // exercise
    bool isHeavy () const { return length > 50; }
    string getHero() const { return hero; }
};

```

Text	title
	author
	length
	topic

Subclasses inherit all members (fields and methods) from their superclasses. All three classes have author, length, methods: getTitle, getAuthor, getLength, isHeavy. ...except that this doesn't work !!!

```

string Text::isHeavy() const { return length > 200; } // private in Book

```

2 options:

```

// 1)
class Book {
    string title, author;
protected:
    int length;
public:
    ...
};

```

```
// 2)
string Text::isHeavy() const { return getLength() > 200; }
```

Recommend: option 2)

- You have no control over what subclasses might do.
- protected weakness encapsulation. Cannot enforce invariants on protected fields.

If you want subclasses to have privileged access

- keep fields private
- provide `protected` `get__` and `set__` methods

Updated object creation/destruction steps:

Created

1. Space is allocated
2. Superclass part is constructed
3. Fields constructed
4. Ctor body

Destroyed

1. Dtor body
2. Fields destructed
3. Superclass is destructed
4. Space deallocated

Must now revisit everything we have learned to see the effect of inheritance.

Type compatibility: Texts and Comics are special kinds of Books.

- should be useable in place of Book.

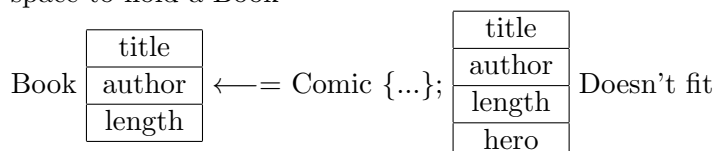
```
Book b = Comic {"", "", 75, ""}; // compiled
b.isHeavy(); // false! 75 is a heavy comic, not a heavy Book.
```

If b is a Comic, why is b acting like a Book?

Because b is a Book. Consequence of stack-allocated objects.

sets aside enough

space to hold a Book



Keeps only the Book part

- comic part is chopped off - "slicing"
- so it really is just a book now.
- so `Book::isHeavy` runs

- slicing happens even if superclass and subclass are the same size.

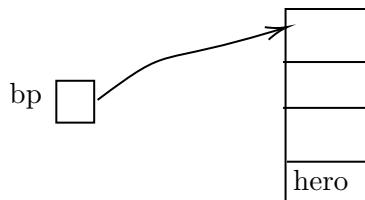
```
vector<Book> library;
library.push_back(Comic{...});
// only the Book parts will be pushed
// not a heterogeneous collection
```

But note:

```
void f(Book books[]); // raw arrays
Comic comics[] = {...};
f(comics);
// will succeed (legal) - but *never* do this!
// array will be misaligned
// will not act like an array of Books
```

Slicing does not happen through ptrs:

```
Book *bp = new Comic {"", "", 75, ""};
```



```
bp->isHeavy(); // still false
```

Rule: the choice of which `isHeavy` to run is based on the type of the pointer (static type), not the object (dynamic type)

Why? Cheaper.

C++ design principle: if you don't use it, you shouldn't have to pay for it. i.e. if you want something expensive, you have to ask for it.

To make `*bp` act like a `Comic` when it is a `Comic`:

```
class Book {
    ...
public:
    ...
    virtual bool isHeavy() const {...}
};
class Comic {
```

```

...
public:
    ...
    bool isHeavy() const override {...}
};
bp->isHeavy(); // true!

```

Assume `isHeavy` is `virtual`.

can now have a true heterogeneous collection:

```

vector<Book*> library;
library.push_back(new Book {...});
library.push_back(new Comic {...});
// even better
vector<unique_ptr<Book>> library;

int howManyHeavy(const vector<Book*> &v) {
    int count = 0;
    for (auto &b : v) {
        if (b->Heavy()) ++count;
        // correct version of isHeavy is always chosen, even though we don't
        // know what's in the vector, and the items are probably not the same type.
    }
    return count;
}

for (auto &b : library) delete b; // Not necessary if you use unique_ptr

```

polymorphism

How do virtual methods "work" and why are they more expensive? Implementation-dependent, but the following is typical: see Figure 18.1

Non-virtual methods \implies ordinary function calls

If there is at least one virtual method

- compiler creates a table of function pointer - one per class - the vtable.
- each object contains a ptr to its class' vtable - the vptr
- calling a virtual method - follow the vptr to the vtable, follow the function ptr to the correct function

vptr is often the "first" field

- so the subclasses object still looks like a super object
- so the program knows where the vptr is.

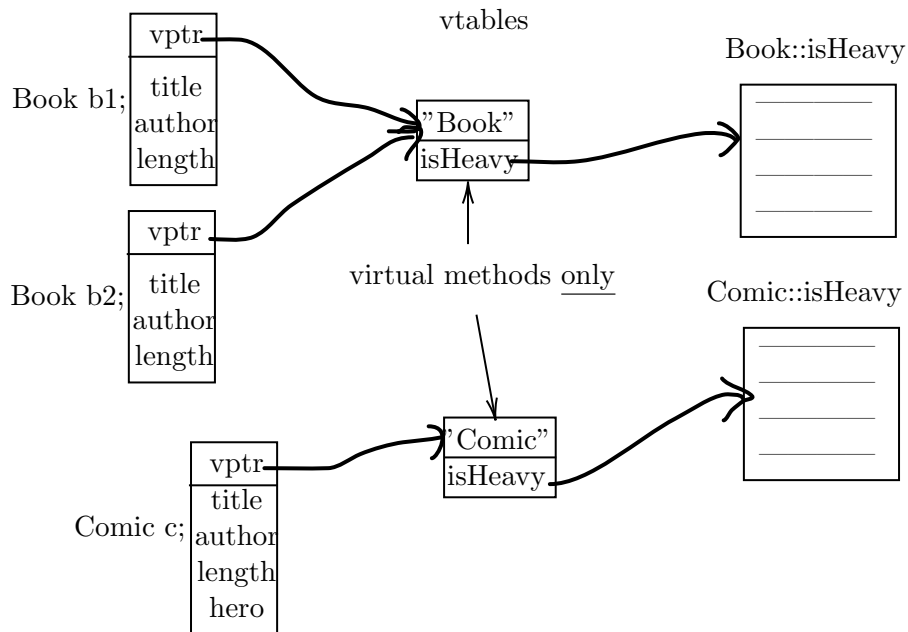


Figure 18.1.: An example of vtable

- so virtual methods incur a cost in time (extra ptr dereferences) and in space (each object gets a `vptr`)
- If a subclass doesn't override a virtual method, its vtable will point at the superclass implementation

Problem 19: I'm Leaking

Consider:

```
class X {
    int *a;
public:
    X (int n): a {new int [n]} {}
    ~X() { delete [] a; }
};

class Y: public X {
    int *b;
public:
    Y(int n, int m): X {n}, b {new int[n]} {}
    ~Y() { delete[] b; }
};
```

Note: Y's dtor will call X's dtor (Step 3)

```
X *px = new Y {3, 4};
delete px; // Leaks! - Calls X's dtor, but not Y's
```

Solution: make the dtor virtual

```
class X {
    ...
public:
    ...
    virtual ~X() {delete [] a;}
};
```

- Always make the dtor virtual in classes that are meant to be superclasses - even if the dtor does nothing. - you never know what the subclass might do, so you need to make sure its dtor gets called.
- Also always give your virtual dtor an implementation - even if empty - it will get called by subclass dtor

If a class is not meant to be superclass

- no need to incur the cost of virtual methods needlessly.
- leave the dtor non-virtual, but declare the class final:

```
class X final {  
    ... // cannot be subclassed  
};
```


Problem 20: I want a class with no objects

```
class Student {
public:
    virtual float fees() const;
};

class RegularStudent: public Student {
public:
    float fees() const override; // Regular student fees
};

class CoOpStudent: public Student {
public:
    float fees() const override; // Co-op student fees
};
```

What should `Student::fees` do?

Don't know - every `Student` should be `Regular` or `Co-op`.

Solution - explicitly give `Student::fees` no implementation.

```
class Student {
public:
    virtual float fees() const = 0; // called a pure-virtual method.
};
```

Abstract classes cannot be instantiated, but can point to instances of create subclasses.

```
// Student s;
// Student *s = new Student;
Student *s = new Regular Student;
```

Subclasses of abstract classes are abstract unless they implement all pure virtual methods.

Abstract classes

- used to organize concrete classes
- can contain common fields, methods, default implementations

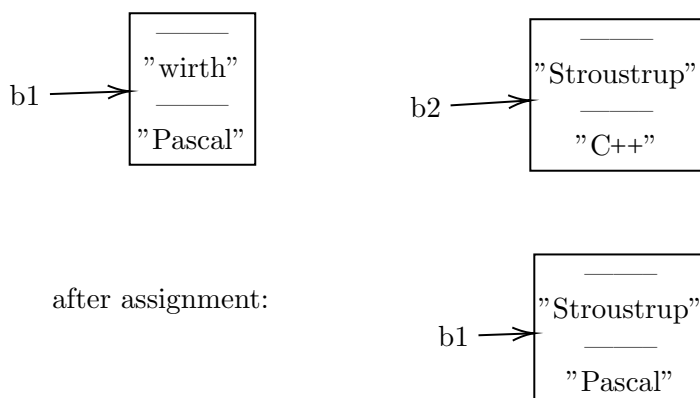
Problem 21: The copier is broken

How do copies and moves interact with inheritance?

```
// copy ctor
Text::Text(const Text &other): Book{other}, topic{other.topic} {}
// move ctor
Text::Text(Text &&other): Book{std::move(other)}, topic{std::move(other.
    topic)} {}
// copy assign
Text &Text::operator=(const Text &other) {
    Book::operator=(other);
    topic = other.topic;
    return *this;
}
// move assign
Text &Text::operator=(Text &&other) {
    Book::operator=(std::move(other));
    topic = std::move(other.topic);
    return *this;
};
```

But consider:

```
Book *b1 = new Text {...}, *b2 = new Text {...};
*b1 = *b2;
```



What happen? Only the Book part is copied - partial assignment

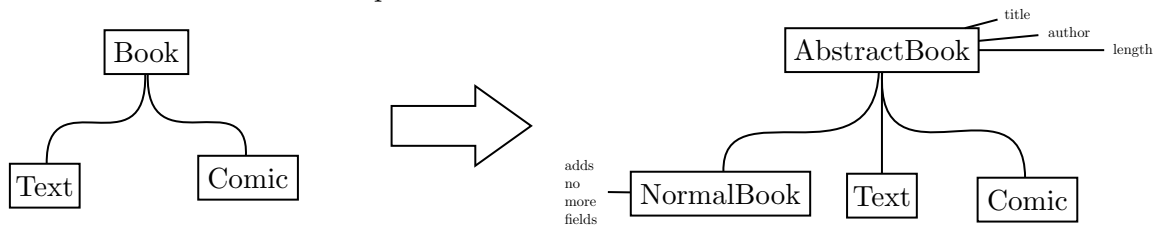
- topic doesn't match title and author - object is corrupted.

Possible solution? Make `operator=` virtual.

```
class Book {
...
public:
    virtual Book &operator=(const Book &other) {...}
};
class Text: public Book {
...
public:
    // Text &operator= (const Text &other) override {...}
    Text &operator= (const Book &other) override {...} // mixed assignment
    // inside () -> must be Book. Or it's not an override.

    // so could pass a Comic - also a problem - will revisit shortly
};
```

Another solution Make all superclasses abstract



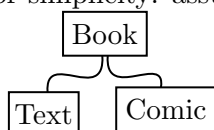
```
class AbstractBook {
protected:
    AbstractBook &operator=(const AbstractBook &other) {...} // non-virtual
    virtual ~AbstractBook() = 0; // to make the class abstract
};
AbstractBook::~AbstractBook() {}

class Text: public AbstractBook {
public:
    Text &operator=(const Text &other) {
        Book::operator=(other);
        topic = other.topic;
    }
};
```

`operator=` non-virtual, so no mixed assignment.
 ∴ no partial assignment, i.e. `*b1=*b2`; won't compile.

Problem 22: I want to know what kind of Book I have

For simplicity: assume



C-style casting:

- (type) expr
- forces 'expr' to be treated as type 'type'

```
int addr = ...;
int *p = (int *) addr;
```

The C++ casting operators: - 4 operators

1. `static_cast` for conversions with a well-defined semantics.

```
void f (int a);
void f (double b);
int x;
f(static_cast<double> (x));
```

superclass ptr to subclass ptr

```
Book *b = new Text {...};
Text *t = static_cast<Text *>(b);
```

You have to know that `b` really points at a `Text` - otherwise undefined behaviour.

2. `reinterpret_cast`

- for cases without a well-defined semantics.
- unsafe, implementation-dependent.

```
Book *b = new Book {...};
int *p = reinterpret_cast<int *>(b);
```

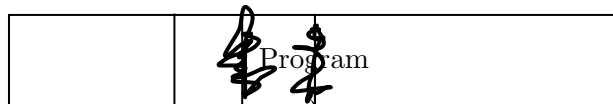
3. `const_cast`

- for adding/removing `const`
- the only C++ `const` that can "cast away `const`"

```
void g (Book &b); // assume g doesn't actually modify b.
void f(const Book &b) {
    g(const_cast<Book&>(b));
}
```

don't use `const`

`const` correct



boundary - where errors occur

4. `dynamic_cast`

```
Book *pb = ...;
```

- what if we don't know whether `pb` points at a `Text`?
- `static_cast` not safe.

```
Text *pt = dynamic_cast<Text *>(pb);
```

If `*pb` is a `Text` or a subclass of `Text`, case succeeds - `pt` points to the object.
else - `pt = nullptr`

```
if (pt) { ... pt->getTopic() ...
else ...// not a Text
```

eg

```
void whatIsIt (Book *pb) {
    if (dynamic_cast<Text *>(pb)) cout << "Text";
    else if (dynamic_cast<Comic *>(pb)) cout << "Comic";
    else cout << "Book";
}
```

- not good style - What happens when you create a new `Book` type?

Dynamic reference casting:

```
Book *pb=...;
Text &t = dynamic_cast<Text &>(*pb);
```

If `*pb` is a `Text` - OK. else raises `std::bad_cast`

Note dynamic casting works by accessing on object's Run-Time Type Information (RTTI)
- this is stored in the vtable for the class.

∴ can only use `dynamic_cast` on objects with at least one virtual method.

Dynamic reference casting offers a possible solution to the polymorphic assignment problem.

```
Text &Text::operator=(const Book &other) { // virtual
    const Text &textother = dynamic_cast<const Text &>(other);
    // throws if other is not a Text
    if (this == &other) return *this;
    Book::operator=(other);
    topic = textother.topic;
    return *this;
}
```

Part II.

A Big Unit on Objected-Oriented Design

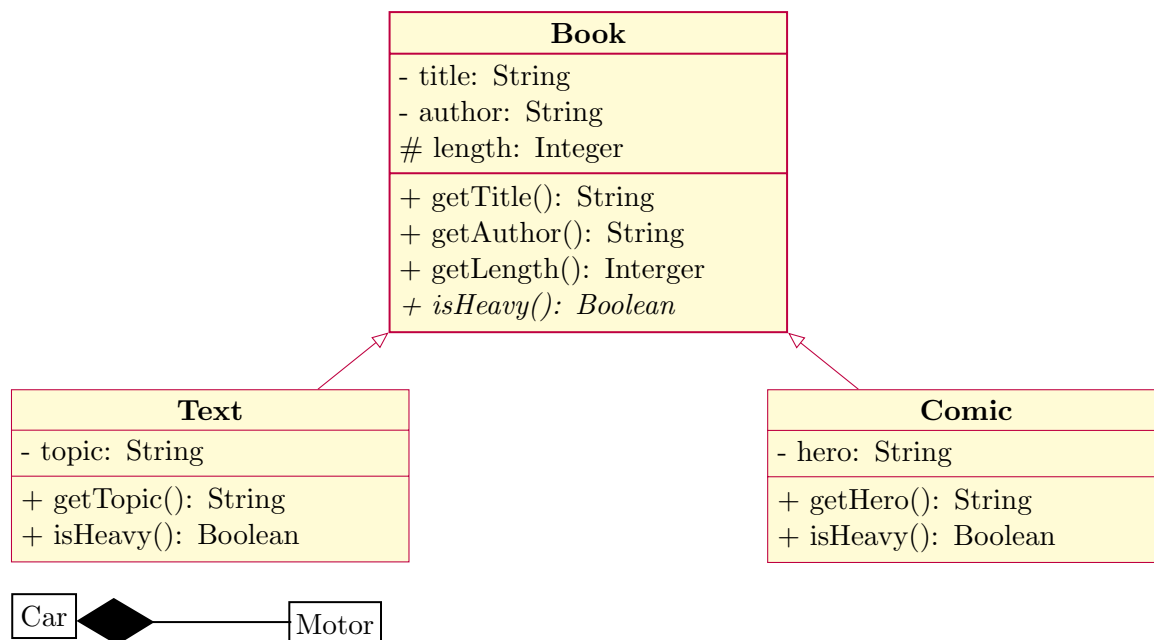
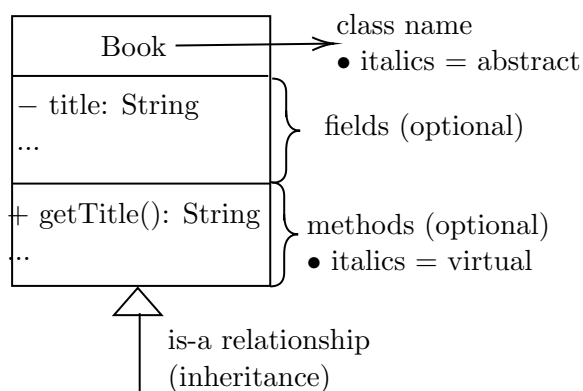
23. UML	88
24. Measures of Design Quality	90
24.1. Coupling & Cohesion	90
25. SOLID Principles of OO design	92
25.1. <u>S</u> ingle Responsible Principle	92
25.2. <u>O</u> pen-Closed Principle	93
25.3. <u>L</u> iskov Substitution Principle	94
25.4. <u>I</u> nterface Segregation Principle	98
25.5. <u>D</u> ependency Inversion Principle	102
26. Some More Design Patterns	105
26.1. Factory Method Pattern	105
26.2. Decorator Pattern	106
26.3. Visitor Pattern	111

System Modelling - UML (United Modelling Language)

- makes ideas easy to communicate, aids design discussion

Notation

- = private
 # = protected
 + = public

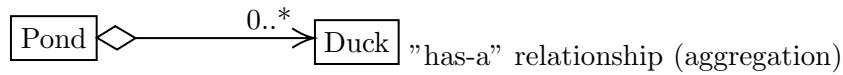


owns "a" relationship

- means the motor is part of the car - does not have an independent existence
- copy/destroy the car \implies copy/destroy the motor

- typical implementation: class composition, i.e. object fields.

```
class Car {Motor m; };
```



- Duck has its own independent existence
- copy/destroy pond \Rightarrow copy/destroy ducks
- typical implementation

```
class Pond { vector<Duck*> ducks; };
```

The concept of ownership is central to OO design in C++ . Every resource (memory, file, window, etc.) should be owned by an object that releases it. RAII.

A `unique_ptr` owns the memory it points to.

A raw ptr should be regarded as not owning the memory it points to.

If you need to point at the same object with several ptrs, one ptr should own it and and be `unique_ptr`. The rest should be raw ptrs.

Moving a `unique_ptr` = transferring ownership

If you need true shared ownership - later.

24.1. Coupling & Cohesion

Coupling - how strongly different modules depend on each other?

LOW:

- function calls with parameters/results of basic type
- function calls with array/struct parameters
- modules that affect each other's control flow

HIGH:

- modules access each other's implementation (friends)

high coupling \implies

- changes to one module affect other modules
- harder to reuse individual modules

Eg function `WhatIsIt` (`dynamic_cast`)

- tightly coupled to the Book class hierarchy
- must change this function if you create another Book subclass

Cohesion - how closely are elements of a module related to each other?

LOW:

- arbitrary grouping (e.g. `<utility>`)
- common theme, otherwise unrelated, maybe some common base code (e.g. `<algorithm>`)

- elements manipulate state over the lifetime of an object (e.g. open/read/close files)
- elements pass data to each other

HIGH:

- elements cooperate to perform exactly one task

low cohesion \implies poorly organized code; hard to understand, maintain

Want low coupling, high cohesion.

SOLID Principles of OO design

25.1. Single Responsible Principle

- A class should only have one reason to change
i.e. a class should do one thing, not several
- Any change to the problem spec requires a change to the program
- If change to ≥ 2 different part of the spec cause changes to the same class, SRP is violated.

Eg Don't let your classes print things.
Consider:

```
class ChessBoard {  
...  
    cout << "Your_move";  
...  
};
```

- bad design - inhibits code reuse.
- What if you want a version of your program that
 - communicates over different streams (files, network)?
 - works in another language?
 - uses graphics instead of that?
- Major changes to ChessBoard class, instead of reuse.

Violates SRP

- must change this class if there is any change to the specification for
 - game rules
 - strategy
 - interface

- etc.
- low cohesion

Split these responsibilities up

- one module (not main! can't reuse main) responsible for communication.
- pass info to the communications object and let it do the talking

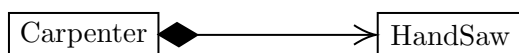
On the other hand

- specifications that are unlikely to change may not need their own class
- avoid needless complexity - judgement call

25.2. Open-Closed Principle

- classes, modules, functions, etc. Should be open for extension and closed for modification.
- Changes in a program's behaviour should happen by writing new code - extending functionality - not changing old code

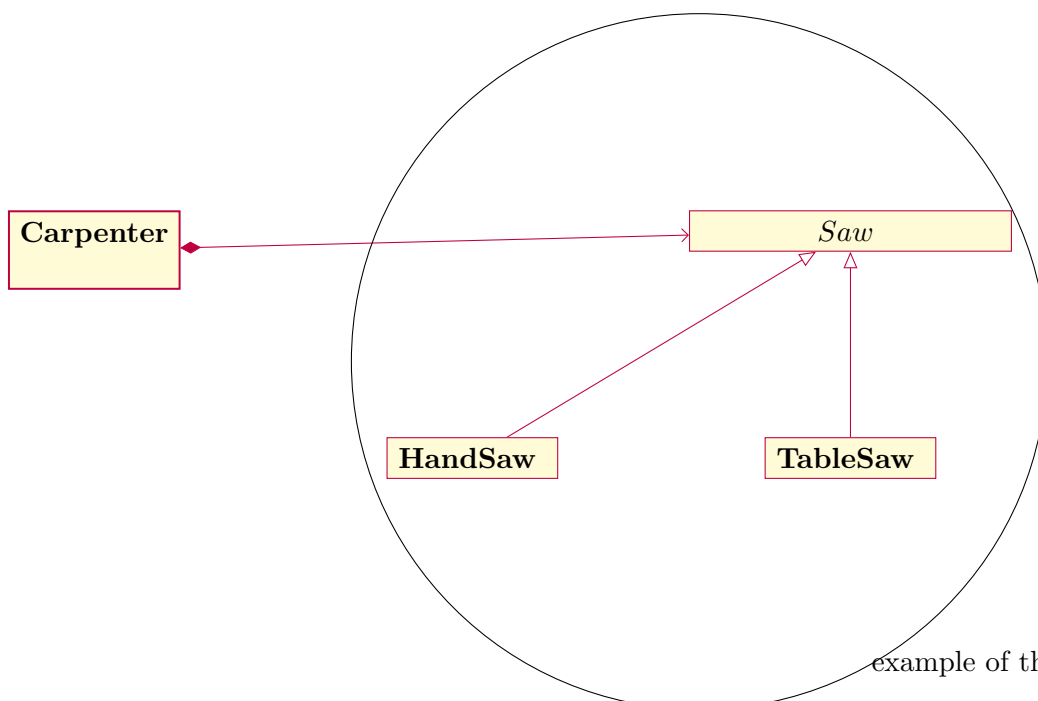
Eg



What if the carpenter buys a tableSaw?

- this design is not open for extension - must change code.

Solution



example of the Strategy Pattern

No changes need for Carpenter.

Also note countHeavy function

```
int countHeavy(const vector<Book *> &v) {
    int count = 0;
    for (auto &p:v) if (p->isHeavy()) ++count;
}
```

- no changes needed if new Book invented.
- vs. WhatIsIt (dynamic casting) - not closed to modification

Note can't really be 100% closed. Some changes require source modifications.
Plan for the most likely changes and make code closed with respect to those changes.

25.3. Liskov Substitution Principle

- Simply put: public inheritance must indicate an "is-a" relationship.

But there is more to it than that:

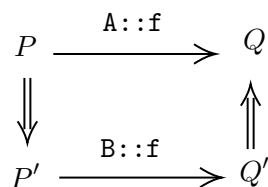
If B is a subtype (subclass) of A, then we should be able to use an object b of type B in any context that requires an object of type A, without affecting the correctness of the program. This is the important point.

C++'s inheritance rules already allow us to use subclass objects in place of superclass objects.

Informally: a program should "not be able to tell", if it is using a superclass object or a subclass object.

More formally

- If an invariant I is true of class A, then it must be true of class B.
- If an invariant I is true of method of $A::f$, and $B::f$ overrides $A::f$, then I must hold for $B::f$
- If $A::f$ has precondition P and postcondition Q , then $B::f$ must have precondition $P' \Leftarrow P$ and postcondition $Q' \Rightarrow Q$.

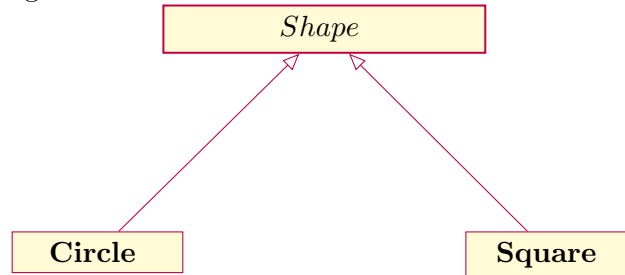


- If $A::f$ and $B::f$ behave differently, the difference in behaviour must fall within what is allowed by the program's correctness specification.

Examples

- contravariance problem
 - arises any time you write a binary operator, i.e. a method with an "other" parameter of the same type as `*this`.

eg



```

class Shape {
public:
    virtual bool operator==(const Shape &other) const;
};
class Circle:public Shape {
public:
    bool operator==(const Circle &other) const override;
};
  
```

- violates Liskov substitution.
 - A Circle is a Shape
 - A Shape can be compared with any other shape.
 \therefore A Circle can be compared with any other shape.
 - (We saw this with `virtual operator=`)
 - C++ will flag this as a compiler error.

Fix

```

#include <typeinfo>
bool Circle::operator==(const Shape &other) const {
    if (typeid(other) != typeid(Circle)) return false;
    const Circle &other= static_cast<const Circle &> (other);
    //compare fields between *this and other
}
  
```

dynamic_cast vs typeid

```

dynamic_cast<const Circle &> (other) // is other a Circle or a subclass
                                     of Circle?
typeid(other) == typeid(Circle) // is other precisely a circle?
// typeid(other) returns an object of type type_info
  
```

- Is a square a rectangle?
 - a square has all the properties of the rectangle.

```

class Rectangle {
    int length, width;
public:
    int getLength() const;
    int getWidth() const;
    virtual void setLength (int length);
    virtual void setWidth (int width);
    int area() const { return length*width; }
};

class Square: public Rectangle {
public:
    Square(int side): Rectangle {side, side} {}
    void setLength(int length) override {
        Rectangle::setLength(length);
        Rectangle::setWidth(length);
    }
    // setWidth - similar
};

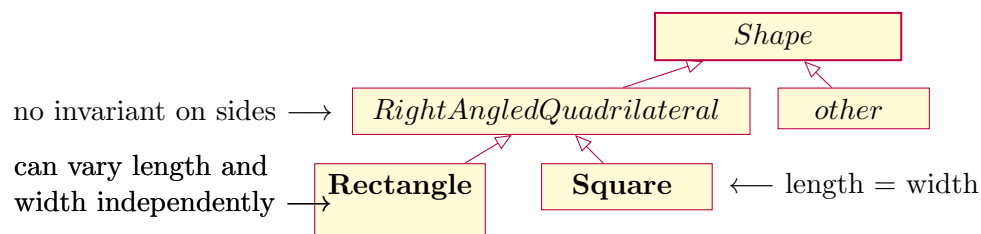
int f(Rectangle &r) {
    r.setLength(10);
    r.setWidth(20);
    return r.area(); // expect 200
}

Square s{1};
f(s); // 400

```

Rectangles have the property that their length and width can vary independently; square do not. So this violates LSP.

On the other hand, an immutable Square could substitute for an immutable rectangle.



Constraining what subclasses can do.
Consider

```

class Turtle {
public:
    virtual void draw = 0;
};

class RedTurtle: public Turtle {
public:
    void draw () override {

```



```

        drawHead();
        drawRedShell();
        drawTail();
    }
};

class GreenTurtle: public Turtle {
public:
    void draw () override {
        drawHead();
        drawGreenShell();
        drawTail();
    }
};
// head and tail are same

```

- code duplication
- Plus - how can we ensure that overrides always do these things?

```

class Turtle {
public:
    void draw () {
        drawHead();
        drawShell();
        drawTail();
    }
private:
    void drawHead();
    virtual void drawShell();
    void drawTail();
};

class RedTurtle: public Turtle {
    void drawShell() override;
};

class GreenTurtle: public Turtle {
    void drawShell() override;
};
// subclass can override a private method, but is not able to call it.

```

Subclasses cannot control the steps of drawing a Turtle, nor the drawing of head & Tail. Can only control the drawing of the shell.
 - called the Template Method Pattern

Extension: NVI (Non-Virtual Interface) idiom

- public: virtual methods are simultaneously
 - part of a class' interface - pre/post conditions class invariants, stated purpose.
 - "hooks" for customization by subclasses - overriding code could be anything. These two at odds with each other.
- hard to vary these independently if they are bound to the same function

NVI says: all virtual methods should be private, i.e. all public methods should be non-virtual.

Eg

```
class DigitalMedia {
public:
    virtual void play() = 0;
};
```

↓

```
class DigitalMedia {
public:
    void play() {
        // here can add before/after code
        // eg: checkCopyright()    updatePlayCount()
        doPlay();
    }
private:
    virtual void doPlay = 0;
};
```

- Generalizes Template Method Pattern - puts every virtual function inside a template method.

25.4. Interface Segregation Principle

- many small interfaces is better than one large interface
- if a class has many functionalities, each client of the class should see only the functionality if needs.

Eg video game

```
class Enemy {
public:
    virtual void draw(); // needed by UI
    virtual void strike(); // needed by game logic
};

class UI {
    vector <Enemy *> v;
```

```
};
class Battlefield {
    vector <Enemy *> v;
};
```

If we need to change the drawing interface, **Battlefield** must recompile for no reason. If we need to change the combat interface, **UI** must recompile for no reason.

Creates needless coupling between **UI** and **Battlefield**

One solution: multiple inheritance

```
class Draw { // actually means drawable
public:
    virtual void draw() = 0;
};
class Combat { // combatable
public:
    virtual void strike() = 0;
};

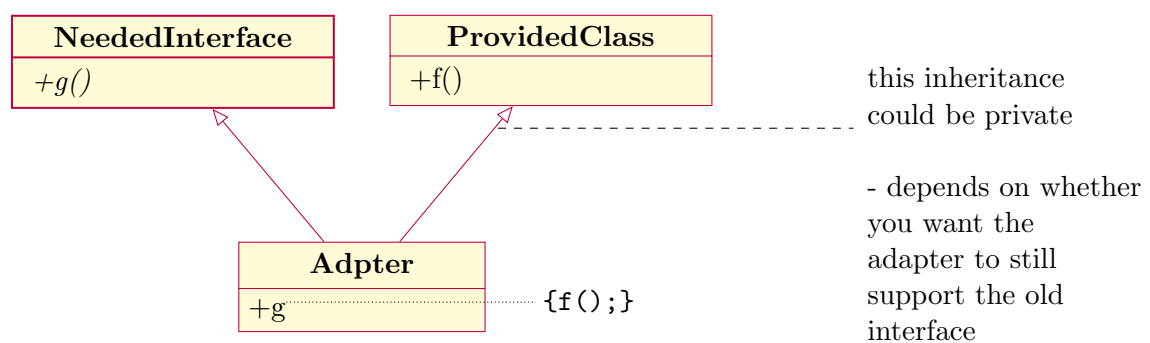
class Enemy: public Draw, public Combat { };

class UI {
    vector <Draw *> v;
};
class Battlefield {
    vector <Combat *> v;
};
```

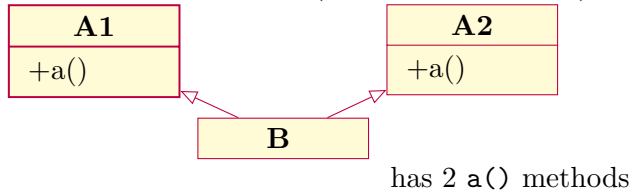
Example of the Adapter Pattern.

General use of Adapter - when a class provides an interface different from the one you need.

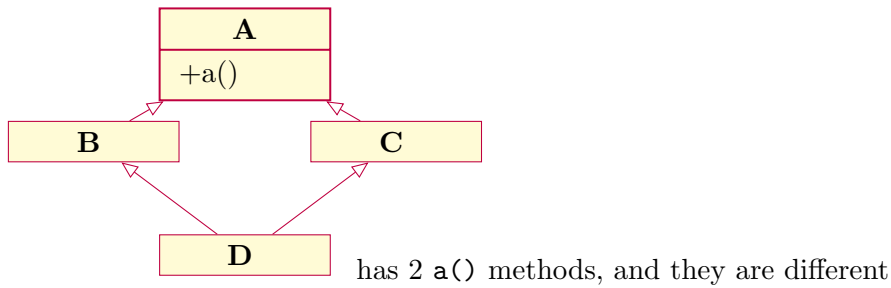
Eg



Detour issue with MI (multiple inheritance)



or indeed



```

class D: public B, public C {
    void f() { ... a() ...} // use B::a() or C::a(). ambiguous
};

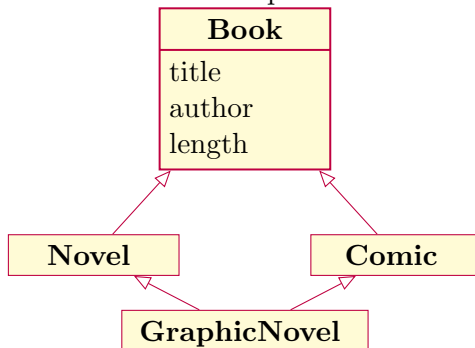
D d;
d.a(); // ambiguous
// use d.B::a() or d.C::a()
  
```

OR may be there should be only one A base and therefore only one a().

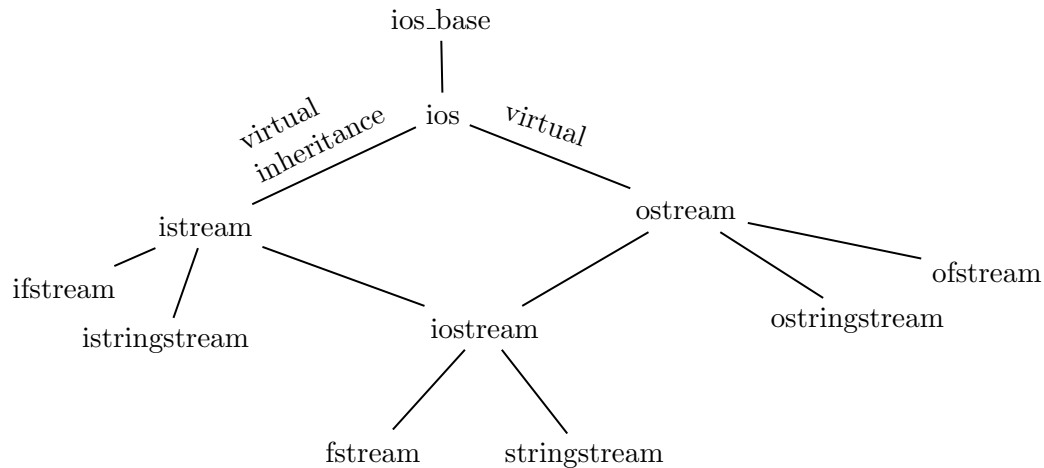
```

class B: virtual public A /* virtual base class */ {...}; // virtual
    inheritance
class C: virtual public A {...};
d.a(); // No longer ambiguous
  
```

Recall Book example:

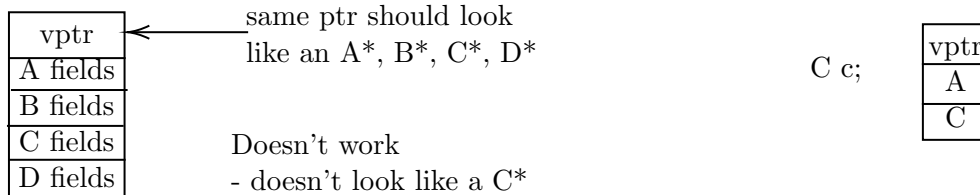


Eg iostream hierarchy

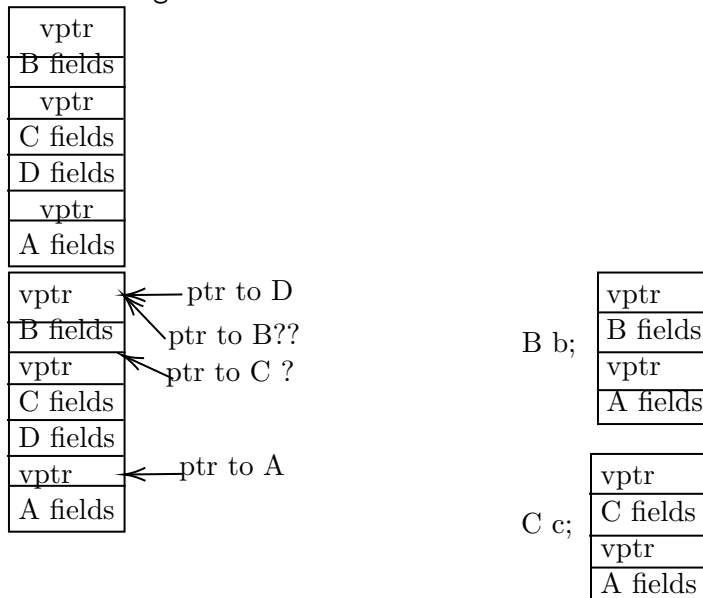


Problem How will a class like D be laid out in memory? (implementation-specific)

Consider



What does g++ do?



B & C need to be laid out so that we can find the A part, but the distance is not known (depends on the run-time type of the object).

Solution: location of the base object stored in vtable.

Also note: diagram doesn't simultaneously look like A, B, C, D.

- but slices of it do.
- \therefore ptr assignment A, B, C, D may change the address stored in the ptr.

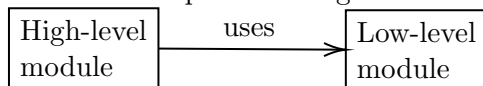
```
D *d = ....;
A *a = d; // changes the address - adds an offset to point to the A
           part.
```

`static_cast` and `dynamic_cast` under MI will also adjust the value of the ptr. `reinterpret_cast` will not.

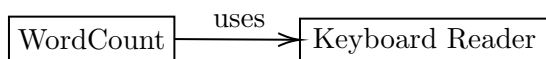
25.5. Dependency Inversion Principle

- High-level modules should not depend on low-level modules. Both should depend on abstractions.
- Abstract classes should never depend on concrete classes.

Traditional top-down design:

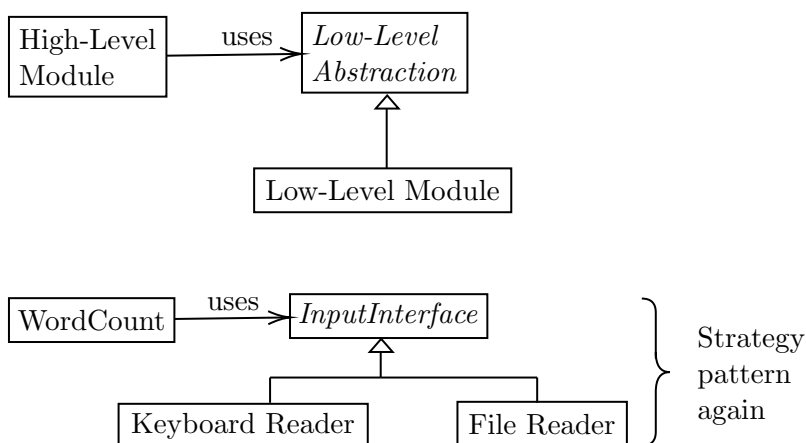


Eg

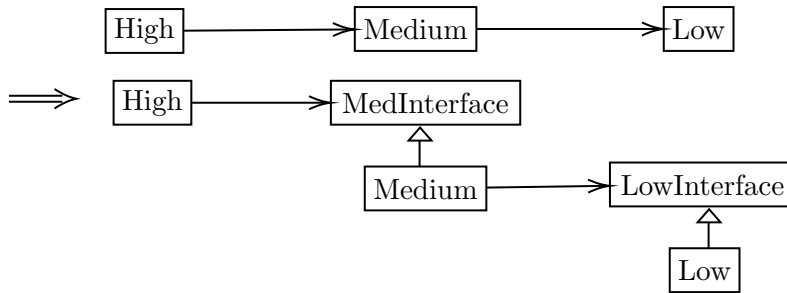


What if I want to use a file reader?
- changes to details affect higher-level module.

Dependency inversion



Works over several layers as well:

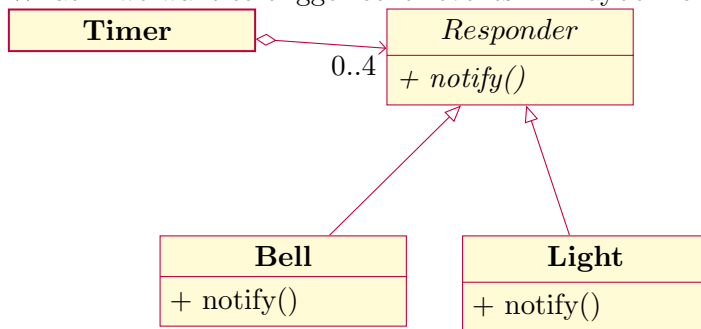


Eg

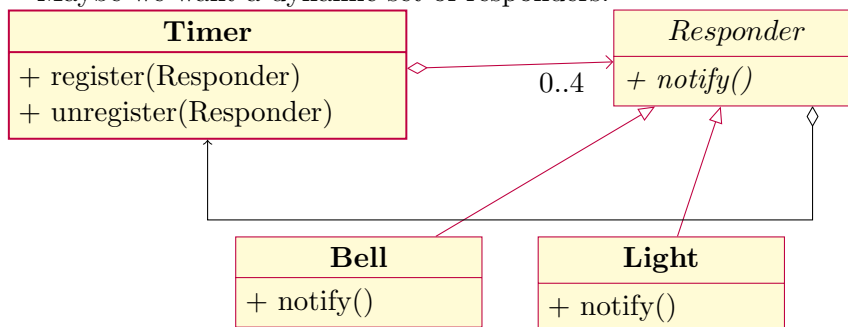


When the timer hits some specified time, it rings the bell (call `Bell::notify`, which rings the bell)

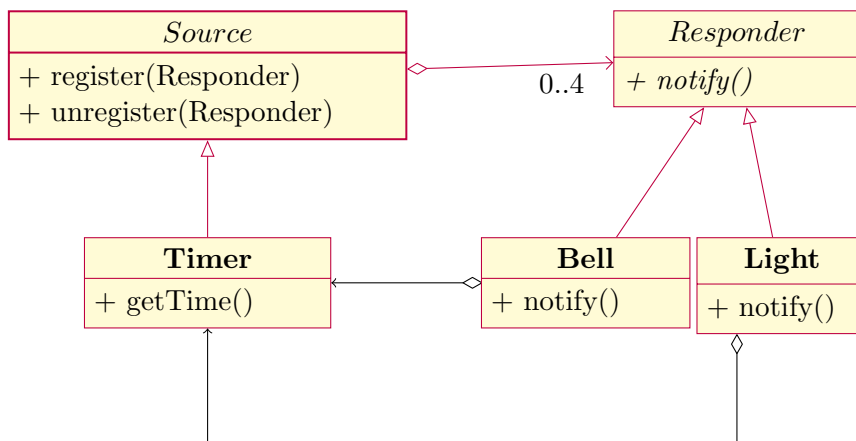
What if we want to trigger other events? - Maybe more than one?



Maybe we want a dynamic set of responders:



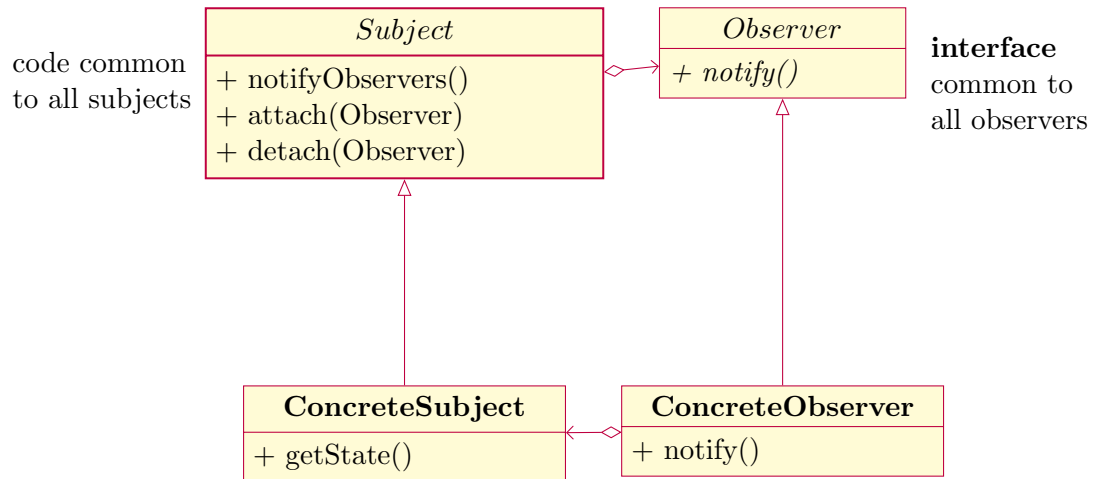
Now Responder is depending on concrete Timer class. Can apply DIP again:



If Light/Bell's behaviour depends on the time, they may need to depend on the concrete timer for a `getTime` method.

- could dependency invert this as well, if you wanted.

General Solution known as Observer Pattern.



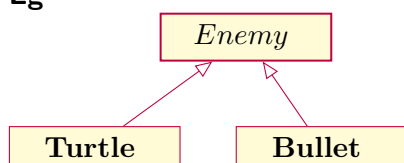
Sequence of calls:

1. Subject's state changes
2. **Subject::notifyObservers** - calls each observer's `notify`
3. Each Observer calls **ConcreteSubject::getState** to query the state and react accordingly.

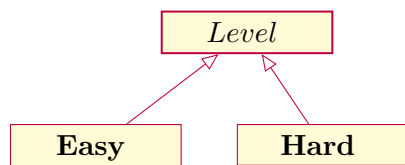
26.1. Factory Method Pattern

- when you don't know exactly what kind of object you want, and your precedences may vary.
- also called the virtual constructor Pattern.
- Strategy Pattern applied to object creation.

Eg



- randomly generated
- more turtles in easy levels
- more bullets in hard levels



```

class Level {
public:
    virtual Enemy *getEnemy() = 0;
};

class Easy: public Level {
public:
    Enemy *getEnemy () override {
        // most turtles
    }
};
  
```

```

class Hard: public Level {
public:
    Enemy *getEnemy () override {
        // most bullets
    }
};

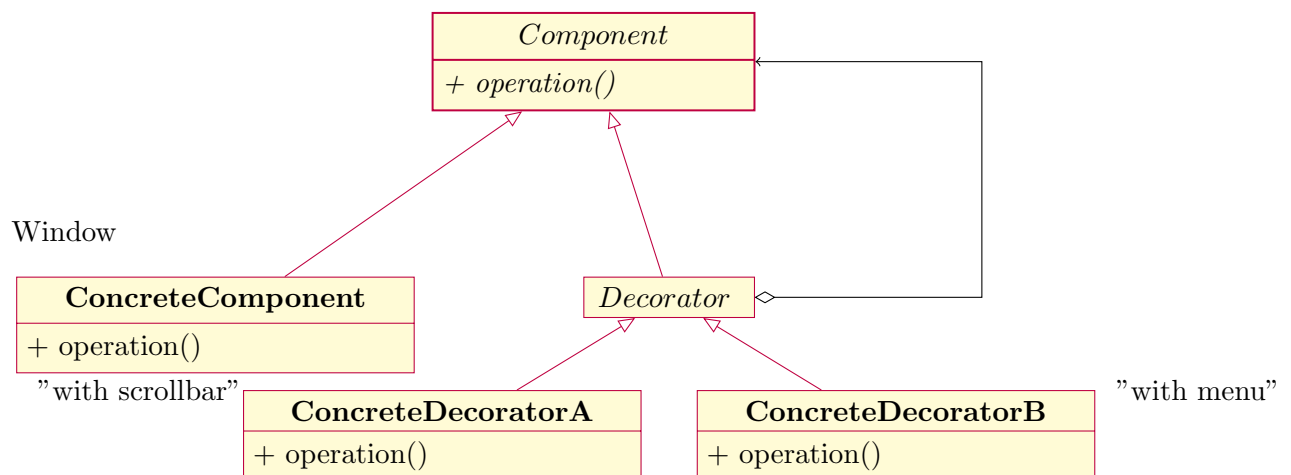
Level *l = ...;
Enemy *e = l->getEnemy();

```

26.2. Decorator Pattern

add/remove functionality to/from objects dynamically.

Eg add menu/scroll bar to basic windows - without a combinatorial explosion of subclasses



- Every a Decorator IS a Component and HAS a Component.
- WindowWithScrollbar IS a Window and has a ptr to the underlying plain window.
- Window w/ Scrollbar & menu is a Window and has a ptr to Window w/ scrollbar, which has a ptr to a Window.

```

WindowInterface *w = new WindowWithMenu{ new WindowWithScrollbar { new
    Window {}}};

```

Example pizza

```

// pizza.h
#ifndef _PIZZA_H_
#define _PIZZA_H_
#include <string>

```

```

class Pizza {
public:
    virtual float price() = 0;
    virtual std::string description() = 0;
    virtual ~Pizza();
};

#endif

// pizza.cc
#include "pizza.h"

Pizza::~Pizza() {}

// decorator.h
#ifndef _DECORATOR_H_
#define _DECORATOR_H_
#include "pizza.h"

class Decorator: public Pizza {
protected:
    Pizza *component;
public:
    Decorator(Pizza *component);
    virtual ~Decorator();
};

#endif

// decorator.cc
#include "pizza.h"
#include "decorator.h"

Decorator::Decorator(Pizza *component): component{component} {}

Decorator::~Decorator() { delete component; }

// crustandsauce.h
#ifndef _CRUSTANDSAUCE_H_
#define _CRUSTANDSAUCE_H_
#include <string>
#include "pizza.h"

class CrustAndSauce: public Pizza {
public:
    float price() override;
    std::string description() override;
};

#endif

```

```

// c & s.cc
#include <string>
#include "crustandsauce.h"
using namespace std;

float CrustAndSauce::price() { return 5.99; }

string CrustAndSauce::description() { return "Pizza"; }

// dippingsauce.h
#ifndef _DIPPINGSAUCE_H_
#define _DIPPINGSAUCE_H_
#include "pizza.h"
#include "decorator.h"
#include <string>

class DippingSauce: public Decorator {
    std::string flavour;
public:
    DippingSauce(std::string flavour, Pizza *component);
    float price() override;
    std::string description() override;
};

#endif

// dipping.cc
#include "dippingsauce.h"
#include <string>
using namespace std;

DippingSauce::DippingSauce(string flavour, Pizza *p):
    Decorator{p}, flavour{flavour} {}

float DippingSauce::price() { return component->price() + .30; }

string DippingSauce::description() {
    return component->description() + " with " + flavour + " dipping sauce";
}

// stuffedcrust.h
#ifndef __STUFFEDCRUST_H__
#define __STUFFEDCRUST_H__
#include "pizza.h"
#include "decorator.h"
#include <string>

class StuffedCrust: public Decorator {
public:
    StuffedCrust(Pizza *component);

```

```

    float price() override;
    std::string description() override;
};

#endif

// stuffedcrust.cc
#include "stuffedcrust.h"
#include <string>
using namespace std;

StuffedCrust::StuffedCrust(Pizza *p): Decorator{p} {}

float StuffedCrust::price() { return component->price() + 2.69; }

string StuffedCrust::description() {
    return component->description() + " with stuffed crust";
}

// topping.h
#ifndef _TOPPING_H_
#define _TOPPING_H_
#include "pizza.h"
#include "decorator.h"
#include <string>

class Topping: public Decorator {
    std::string theTopping;
    const float thePrice;
public:
    Topping(std::string topping, Pizza *component);
    float price() override;
    std::string description() override;
};

#endif

// topping.cc
#include "topping.h"
#include <string>
using namespace std;

Topping::Topping(string topping, Pizza *component):
    Decorator{component}, theTopping{topping}, thePrice{0.75} {}

float Topping::price() { return component->price() + thePrice; }

string Topping::description() {
    return component->description() + " with " + theTopping;
}

```

```

// main.cc
#include <iostream>
#include <string>
#include <vector>
using namespace std;

#include "pizza.h"
#include "topping.h"
#include "stuffedcrust.h"
#include "dippingsauce.h"
#include "crustandsauce.h"

int main() {
    Pizza *myPizzaOrder[3];
    myPizzaOrder[0] = new Topping{"pepperoni",
                                   new Topping{"cheese", new CrustAndSauce}};
    myPizzaOrder[1] = new StuffedCrust{
        new Topping{"cheese",
                    new Topping{"mushrooms",
                                new CrustAndSauce}}};
    myPizzaOrder[2] = new DippingSauce{"garlic",
                                        new Topping{"cheese",
                                                    new Topping{"cheese",
                                                                new Topping{"cheese",
                                                                    new CrustAndSauce}}}}};

    float total = 0.0;

    for (int i = 0; i < 3; ++i) {
        cout << myPizzaOrder[i]->description()
              << ": " << myPizzaOrder[i]->price() << endl;
        total += myPizzaOrder[i]->price();
    }

    cout << endl << "Total cost: " << total << endl;

    for (int i = 0; i < 3; ++i) {
        delete myPizzaOrder[i];
    }
}

```

Makefile

```

CXX = g++-5
CXXFLAGS = -std=c++14 -Wall -MMD
EXEC = main
OBJECTS = main.o pizza.o crustandsauce.o topping.o dippingsauce.o
          stuffedcrust.o decorator.o

```

```

DEPENDS = ${OBJECTS:.o=.d}

${EXEC}: ${OBJECTS}
    ${CXX} ${CXXFLAGS} ${OBJECTS} -o ${EXEC}

-include ${DEPENDS}

.PHONY: clean

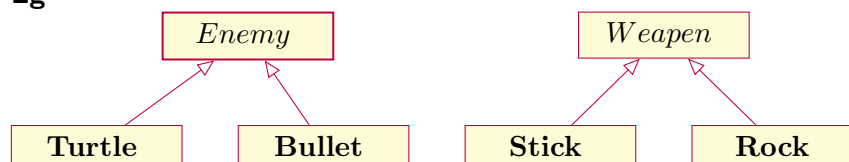
clean:
    rm ${OBJECTS} ${EXEC} ${DEPENDS}

```

26.3. Visitor Pattern

- for implementing double dispatch
- method chosen based on the runtime type of 2 objects, rather than just one.

Eg



- effect of striking an enemy with a weapon depends on both the enemy and the weapon
- C++
 - virtual methods are dispatched on the type of the receiver object, and not the method parameters.
 - No way to specify 2 receiver objects.

Visitor Pattern - combine overriding with overloading to do 2-stage dispatch

```

class Enemy {
public:
    virtual void beStruckBy(Weapon &w)=0;
};

class Turtle: public Enemy {
    void beStruckBy(Weapon &w) override {w.strike(*this);}
};

class Bullet: public Enemy {
    void beStruckBy(Weapon &w) override {w.strike(*this);}
}; // *this has a different type in each body

class Weapon {
public:
    virtual void strike(Turtle &t)=0;
    virtual void strike(Bullet &b)=0;
};

```

```
};

class Stick: public Weapon {
public:
    virtual void strike(Turtle &t) { /* strike Turtle with Stick */ }
    virtual void strike(Bullet &b) { /* strike Turtle with Bullet */ }
};
// etc

Enemy *e = new Bullet {...};
Weapon *e = new Rock {...};
e->beStruckBy(*w); // What happens?
```

Bullet::beStruckBy runs (virtual method dispatch)

- calls Weapon::strike(Bullet &) (known at compile time - overload resolution) (*this is Bullet)
- virtual method resolves to Rock::strike(Bullet &)

Visitor can also be used to add functionality to a class hierarchy without adding new virtual methods.

Add a visitor to the Book hierarchy.

```
class Book {
public:
    virtual void accept(Bookvisitor &v) {v.visit(*this);}
    ...
};

class Text: public Book {
public:
    void accept (BookVisitor &v) override { v.visit(*this); }
};
// etc

class BookVisitor {
public:
    virtual void visit (Book &b) = 0;
    virtual void visit (Text &t) = 0;
    virtual void visit (Comic &c) = 0;
};
```

	For	by
Eg Categorize & count:	Books	Author
	Texts	Topic
	Comics	Hero

Could add a virtual method to the Book hierarchy. OR write a visitor:

```
class Catalogue: public BookVisitor {
public:
```



```

map<string, int> theCat;
void visit(Book &b) override{ ++theCat[b.getAuthor()]; }
void visit(Book &b) override{ ++theCat[b.getTopic()]; }
void visit(Book &b) override{ ++theCat[b.getHero()]; }
};

```

But it won't compile!

- Circular include dependencies book.h, BookVisitor.h include each other
- include guard prevents multiple inclusion
- Whichever ends up occurring first will refer to things not yet defined.

Know when on `#include` is actually needed!

Needless `#includes` create artificial compilation dependencies and slow down compilation
- or prevent compilation all together.

Sometimes a forward class declaration is good enough.

Consider

```

class A {...}; // A.h

class B { A a; };
class C { A *a };
class D: public A { ... };
class E { A f(A); };
class F {
    A f(A a) { a.someMethod; ...}
};
class G {
    t<A> x; // When t is a template
};

```

Which need includes: B, D, F
forward declaration is OK: C, E
G - depends on how t uses A - should collapse to one of the other cases

Now that class F only needs an include because method f's implementation is present and use a method of A

- a good reason to keep implementation in .cc file
- where possible: forward declare in .h, include in .cc

Notice: B needs to include; C does not.

If we want to break the compilation dependency of B on A, we could translate to C.
More generally:

```
class A1 {}; class A2 {}; class A3 {};
class B {
    A1 a1;
    A2 a2;
    A3 a3;
}; // compilation dependency #include...
```

⇒

```
// b.h
class BImpl; // forward declaration only
class B {
    unique_ptr<BImpl> pImpl; // ptr to implementation
};

// bimpl.h
struct Bimpl {
    A1 a1;
    A2 a2;
    A3 a3;
};
```

b.cc: methods reference pImpl→a1 a2 a3

- b.h is no longer compilation-dependent on a1.h, ...
- called the pImpl idiom

-
- Another advantage of pImpl - ptrs have a non-throwing swap
 - can provide the strong guarantee on a B method by
 - * copying the impl into a new Bimpl structure (heap-allocated) - method modifies the copy
 - * if anything throws, discard the new structure (easy & automatic with `unique_ptr`s)
 - * if all succeeds, swap impl struct ptrs (nothrow)
 - * previous impl automatically destroyed by smart ptr

```
class B {
    unique_ptr<Bimpl>pImpl;
    ...
    void f() {
        auto temp = make_unique<Bimpl>(* pImpl);
        temp->doSth();
        temp->doSthElse();
        std::swap(pImpl, temp); // nothrow
    } // Strong guarantee
};
```

Back to abstraction in C++

Part III.

Abstraction in C++

27. Problem 23: Shared Ownership	116
28. Problem 24: Abstraction over Iterators	119
28.1. Template Metaprogramming	123
29. Problem 25: I want an even faster vector	124
29.1. Move/Forward Implementation	127
30. Problem 26: Collecting Stats	130
31. Problem 27: Resolving Method Overrides At Compile-Time	132
32. Problem 28: Polymorphic Cloning	134
33. Problem 29: Logging	136
33.1. Mixin inheritance	137
34. Problem 30: Total Control	138
35. Problem 31: Total Control ove Vectors & Lists	141
36. Problem 32: A fixed-size allocator	142
37. Problem 34: I want a (tiny bit) smaller vector class	146
37.1. Something Missing from Last Year	146

Problem 23: Shared Ownership

Is C++ hard? No (if you are a client programmer)

But explicit memory management ...

- use `vector` when you need an array
- use `unique_ptr` when you need a heap object
- use stack-allocated objects as much as possible

If you followed these, you should never have to say `delete` or `delete[]`.

But `unique_ptr`s don't respect is-a.

```
unique_ptr<Base> p = unique_ptr<Derived> {...}; // OK
p->virtual_fn(); // runs derived version. Ok.
```

But

```
unique_ptr<Derived> q = ...;
unique_ptr<Base> p = std::move(q); // X
// type error - no conversion between unique_ptr<Derived> & unique_ptr<Base>
```

Easy to fix

```
template <typename T> class unique_ptr {
    ...
    T *p;
public:
    ...
    template <typename u> unique_ptr (unique_ptr<u> &&q): p{q.p} {q.p=nullptr}
    }
    template <typename u> unique_ptr &operator=(unique_ptr<u> &&q) {
        std::swap(p,q.p);
        return *this;
    }
};
```

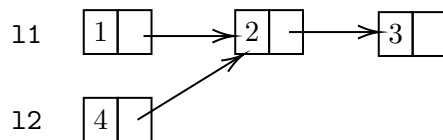
```
// works for any unique_ptr whose pointer is assignment compatible with this
// ->p.
// e.g. subtypes of T, but not supertypes of T
```

But I want two smart ptrs printing at the same object!
Why?

- the ptr that owns the object should be a `unique_ptr`
- all others can be raw ptrs

When would you want true shared ownership?

```
(define l1 (cons 1 (cons 2 (cons 3 empty))))
(define l2 (cons 4 (rest l1)))
```

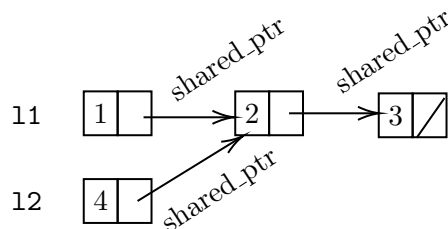


Shared data structures are a nightmare in C.

- How can we ensure each node is freed exactly once?
- Easy in garbage-collected languages.

What can C++ do?

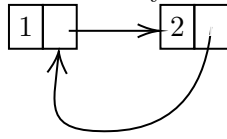
```
template <typename T> class shared_ptr {
    T *p;
    int *refcount; // refcount
    // - counts how many shared_ptrs point to *p
    // - updated each time a refcount is initialized, assigned, or destroyed
    // refcount is shared among all shared_ptrs that point to *p
    // p is only deleted when refcount reaches 0
    // implementation details - left to us
public:
    ...
};
```



```
struct Node {
    int data;
    shared_ptr<Node> next;
};
```

Now deallocation is easy - as easy as garbage collection.

Just watch: cycles



weak_ptrs)

If you have cyclic data - may have to physically break the cycle (or use

Also watch:

```
Book *p = new ...;
shared_ptr<Book> p1 {p};
shared_ptr<Book> p2 {p};
```

- `shared_ptrs` are not mind-readers
- `p1` & `p2` will not share a refcount. BAD
- if you want to point 2 `shared_ptrs` at an object, create one `shared_ptr` and copy it.

But - you can't dynamic cast `shared_ptrs`.

```
template<typename T, typename U> shared_ptr<T> dynamic_pointer_cast <U> &spu
) {
    return shared_ptr<T> (dynamic_cast <T*>(spu.get()));
}

// similarly - const_pointer_cast, static_pointer_cast
```

Just to follow up on the `dynamic_pointer_cast` implementation that Lushman gave us as above. As was observed in class, this implementation doesn't share ref counts with other `shared_ptrs` to the same object. It can be fixed, but it requires knowing about an additional constructor that `shared_ptr` provides:

```
if (dynamic_cast<T*>(spu.get()) return shared_ptr<T>{spu, static_cast<T*>(
    spu.get())};
else return shared_ptr<T>{nullptr};
```

The constructor that was used here basically says, "share ref counts with this `shared_ptr`, but point here instead".

Problem 24: Abstraction over Iterators

I want to jump ahead n spots in any iterator.

```
template <typename Iter> Iter advance(Iter it, size_t n);
```

How should we do it?

```
for (size_t i = 0; i < n; ++i) ++it;
return it;
```

Slow - $O(n)$ time.

Can't just say `it+=n`?

Depends

- for vectors, yes - $O(1)$ time
- for lists, no - `+=` not supported (and if it was, it would still be a loop)

Related - can we go backwards?

- vectors: yes
- lists: no

So all iterators support `!=`, `*` `++`. - but some iterators support other operations.

- `list::iterator` - called a forward iterator - can only go one step forward.
- `vector::iterator` - called a random access iterator - can go anywhere (supports arbitrary ptr arithmetic)

How can we write `advance` to use `+=` for random access iterators, and loop for forward iterators?

Since we have different kinds of iterators, create a class hierarchy.

```
struct input_iterator_tag {};
struct output_iterator_tag {};
struct forward_iterator_tag: public input_iterator_tag {};
struct bidirectional_iterator_tag: public forward_iterator_tag {};
struct random_access_iterator_tag: public bidirectional_iterator_tag {};
```

To associate each iterator class with a tag - could use inheritance.

eg

```
class List {
...
public:
    class iterator: public forward_iterator_tag {...};
};
```

But

- makes it hard to ask what kind of iterator we use (can't `dynamic_cast` - no vtables!)
- doesn't work for iterators that aren't classes (e.g. ptrs)

Instead - make the tag a member:

```
class List {
...
public:
    class iterator {
...
    public:
        typedef forward_iterator_tag iterator_category;
        // convention - every iterator class will define a type member called
        iterator_category.
        ...
    };
};
```

- also doesn't work for iterators that aren't classes
- but we aren't done yet
- make a template that associates every iterator type with its category.

```
template<typename It> struct iterator_traits {
    typedef It::iterator_category iterator_category;
};
```


won't compile for the same reason as on page 70. Instead, we have

```
template<typename It> struct iterator_traits {
    typedef typename It::iterator_category iterator_category;
};
// (typename) What is this doing here?
```

needed so that C++ know that `It::iterator_category` is a type (remember: compiler doesn't know anything about `It`).

Consider:

```
template <typename T> void f() {
    T::something x; // only makes sense if T::something is a type
}

// But
template <typename T> void f() {
    T::something *x;
    // is this a ptr or a multiplication?
    // Can't know unless you whether T::something is a type.
    // C++ assumes it's a value unless told otherwise.
}

template <typename T> void f() {
    typename T::something x;
    typename T::something *y;
}
```

Need to say `typename` whenever you refer to a member type of a template parameter.

(eg

```
iterator_traits<List<T>::iterator>::iterator_category =>
    forward_iterator_tag
```

)

Specialized version for ptrs:

```
template <typename T> struct iterator_traits<T*> {
    typedef random_access_iterator_tag iterator_category;
};
```

For any iterator type `T`, `iterator_traits<T>::iterator_category` resolves to the tag struct type for `T` (including if `T` is a ptr)

What do we do with this?

```
template <typename Iter> Iter advance (Iter it, int n) {
```

```

    if (typeid(typename iterator_traits<Iter>::iterator_category) == typeid(
        random_access_iterator_tag)) {
        return it+=n;
    }
    else if ( ... )
        ...
}
// Doesn't compile.

```

If the iterator is not random access, and \therefore doesn't have a $+=$ operator, $it+=n$ will cause a compilation error, even though it will never be used.

Moreover - the choice of which implementation to use is being made to run-time, when the right choice is known at compile-time.

To make a compile-time decision: overloading

```

template <typename Iter>
Iter doAdvance(Iter it, int n, random_access_iterator_tag) {
    return it += n;
}

template <typename Iter>
Iter doAdvance(Iter it, int n, bidirectional_iterator_tag) {
    if (n > 0) for (int i = 0; i < n; ++i) ++it;
    else if (n < 0) for (int i = 0; i < -n; ++i) --it;
    return it;
}

template <typename Iter>
Iter doAdvance(Iter it, int n, forward_iterator_tag) {
    if (n >= 0) {
        for (int i = 0; i < n; ++i) ++it;
        return it;
    }
    throw SomeError();
}

template <typename Iter> Iter advance (Iter it, int n) {
    return doAdvance(it, n, iterator_traits<Iter>::iterator_category{});
}

```

Now the compiler will select the fast `doAdvance` for random access iterators, the slow `doAdvance` for bidirectional iterators, and the throwing `doAdvance` for forward iterators.

These choices made at compile-time - no runtime cost.

Using template instantiations to perform compile-time computation.
- called template metaprogramming.

28.1. Template Metaprogramming

C++ templates form a functional language that operates at the level of types.

Express conditions via overloading, repetition via recursive template instantiation.

```
template <int N> struct Fact {  
    static const int value = N*Fact<N-1>::value;  
};  
template <> struct Fact <0> {  
    static const int value = 1;  
};  
  
int x = Fact <5>::value; // 120 - evaluated at compile-time
```

But for compile-time computation of values, C++11/14 offer a more straightforward facility.

```
constexpr int fact (int n) {  
    if (n==0) return 1;  
    else return n*fact(n-1);  
}
```

constexpr functions:

- evaluate the function at compile-time if n is a compile-time constant
- else evaluate at run-time
- A constexpr function must be something that actually can be evaluated at compile-time.
 - can't be virtual, can't mutate non-local variables.

Let's solve a problem.

Problem 25: I want an even faster vector

In the good old days of *C*, you could copy an array (even an array of struct!) very quickly by calling `memcpy` (like `strcpy`, but for arbitrary memory, not just strings).

`memcpy` was probably written in assembly, and was as fast as the machine could possibly be.

In C++ - copies invoke copy ctors, which are function calls.

In C++, a type is considered POD (plain old data) if it

- has a trivial default ctor (equiv. `=default`)
- is trivially copyable - copies/moves, dtor have default implementations

AND is standard layout

- no virtual methods or bases
- all members have the same visibility
- no reference members
- no fields in both base class & subclass, or in multiple base classes

For POD types, semantics compatible with C, and `memcpy` is safe to use.

How?

```
template <typename T> class vector {
    T *theVector;
    size_t n, cap;
public:
    ...
    vector (const vector &other): ... {
        if (std::is_pod<T>::value) {
            memcpy(theVector, other.theVector, n*sizeof(T));
        }
        else { //loop
        }
    }
};
```

Works, but ... condition is evaluated at run-time, but result is known at compile-time.
 2nd option:

```
template <typename T> class Vector {
...
public:
    template <typename X=T> vector (enable_if<std::is_pod<X>::value, const T
    &>::type other): ... {
        memcpy(theVector, other.theVector, n*sizeof(T));
    }
    template <typename X=T> vector (enable_if<!std::is_pod<X>::value, const T
    &>::type other): ... {
        // placement new loop
    }
};
```

How does this work?

```
template<bool b, typename T> struct enable_if;
template<typename T> struct enable_if<true, T> {
    using type=T;
};
```

With metaprogramming, what don't you say is as important as what you do say.

If b is true, `enable_if` defines a struct whose 'type' member typedef is T.

So if `std::is_pod<X>::value = true`, then `enable_if<std::is_pod<x>::value, const T&>::type` \Rightarrow `const T&`.

If b is false, then the struct is declared, but not defined.

So `enable_if<b, T>` won't compile.

So one of these two versions of the copy ctor (the one with the false condition) won't compile.

Then how is this a valid program?

C++ rule SFINAE - Substitution Failure Is Not An Error

In other words: if t is a type, `template <typename T> .. f (...) {...}` is a template function, and substituting `T=t` resulting in an invalid function, the compiler does not signal an error - it just removes that instantiation from consideration during overload resolution.

On the other hand - if no version of the function is in scope to handle the overload call, that is an error.

Q Why is this wrong:

```
template <typename T> class vector {
...
public:
    vector(typename enable_if<std::is_pod<T>::value, const T&>::type other):
        ... { ... }
```

```
vector(typename enable_if<!std::is_pod<T>::value, const T&>::type other):
    ... { ... }
};
```

i.e. why do we need the extra template out front?

Because SFINAE applies to template functions, and these methods are ordinary functions (ctors), not templates.

- they depend on T, but T's value was determined when you decided what to put in the vector
- if substituting $T=t$ fails, it invalidates the entire vector class, not just the method.

So make the method a separate template, with a new arg X, which can be defaulted to Tm and do `is_pod<X>`, not `is_pod<T>`.

So it compiles, but when you run it, it crushes.

Why? Hint: if you put debug statement in these copy ctors, they don't print.

Ans: we are getting the compiler-supplied copy ctor, which is doing shallow copies.

- These templates are not enough to suppress the auto-generated copy ctor.
 - and a non-templated match is always preferred to a templated one.

What do we do about it? Could try:

```
template <typename T> class Vector {
...
public:
    vector(const vector &other)=delete; // i.e. disable the copy ctor
    ...
}; // not allowed - can't disable the copy ctor + then write one.
```

Solution that works: overloading.¹

```
template <typename T> class vector {
...
    struct dummy {};
public:
    vector (const vector &other): vector{other, dummy{}} {} // ctor
    delegation
    template<typename X=T> vector(typename enable_if<std::is_pod<T>::value,
        const vector <T>&>::type other, dummy): ... { ... }
    template<typename X=T> vector(typename enable_if<!std::is_pod<T>::value,
        const vector <T>&>::type other, dummy): ... { ... }
};
```

- overload the ctor with an unused "dummy" arg.

¹The solution below corrected the typos above: missing type vector

- have the copy ctor delegate to the overloaded ctor.
- copy ctor is inline, so no function call overhead.

this works

Can write some "helper" definitions to make `is_pod` and `enable_if` easier to use

```
template <typename T> constexpr bool is_pod_v = std::is_pod<T>::value;
template <bool b, typename T> using enable_if_t = typename enable_if<b, T>::
    type;
////////////////////////////////////
template <typename T> class vector {
..
public:
..
    template<typename X=T> vector (enable_if_t <is_pod_v<X>, const vector <T
    >& other, dummy): ... {...}
    template<typename X=T> vector (enable_if_t <!is_pod_v<X>, const vector <T
    >& other, dummy): ... {...}
};
```

We now have the machinery to implement `std::move` and `std::forward`.

29.1. Move/Forward Implementation

`std::move` - first attempt

```
template <typename T> T&& move(T &&x) {
    return static_cast<T &&> (x);
}
```

- doesn't quite work - `T&&` is a universal reference, not an rvalue ref.
If `T` was an lvalue ref, `T&&` will be an lvalue ref.
- Need to make sure that `T` is not an lvalue ref - if `T` is an lvalue ref, get rid of the ref.

`std::move` - correct version

```
template <typename T> inline constexpr std::remove_reference<T>::type &&
    move(T && x) {
    return static_cast<std::remove_reference<T>::type &&>(x);
    // std::remove_reference<T>::type      turns t&, t&& into t
}
```

Q can we save typing and use `auto`?

```
template <typename T> auto move(T &&x) { ... }
```

No! By-value `auto` throws away refs and other consts.

```
int z;
int &y = z;
auto x = y; // x is int

const int &w = z;
auto v = w; // v is int
```

- uses the type a value would have in exprs if copied.

By-ref `auto&&` is a universal reference.

- need a type declaration rule that doesn't discard refs.

- `decltype(...)` - returns the type that ... was declared to have
- `decltype(var)` - returns the declared type of the var
- `decltype(expr)` - returns an lvalue or rvalue ref, depending on whether `expr` is an lvalue or an rvalue

Eg

```
int z;
int &y = z;
decltype(y)x = z; // x is int&
x = 4; // affects z

decltype(z) s=z; // s is int
s = 5; // doesn't affect z

auto x = z; // x is int
x = 4; // does not affect z

decltype((z)) s = z; // s is int &
s = 6; // affects z
```

`decltype(auto)`

- perform type deduction like `auto`, but use the `decltype` rules.
- i.e. don't throw away ref.

Correct move:


```
template <typename T> inline constexpr decltype(auto) move (T &&x) {
    // as before
}
```

std::forward -first attempt

```
template <typename T> T && forward (T &&x) {
    return static_cast<T &&>(x);
}
```

/ Reasoning - if x is an lvalue/rvalue, T&& is an lvalue/rvalue reference.

Doesn't work - forward is called on expressions that are lvalues, but may point at rvalues.

eg

```
template <typename T> void f (T &&y) { // lvalue/rvalue ref
    ... forward(y) ... // y is an lvalue
}
```

In order to work, forward must know what types was deduced for T.

So forward<T>(y) would work.

So - can we prevent automatic type deduction for forward?

```
template <typename T> struct identity {
    using type = T;
};
template <typename T> T &&forward (identity<T>::type &&x) { ... }
// identity<T>::type makes type deduction impossible. but this ref is no
// longer universal.
```

instead - separate lvalue/rvalue cases.

```
template <typename T> inline constexpr T & forward(std::remove_reference_t<T>
    & x) noexcept {
    return static_cast<T &>(x);
}
template <typename T> inline constexpr T && forward(std::remove_reference_t<
    T> &&x) noexcept {
    return static_cast<T &&>(x);
}
```

Problem 26: Collecting Stats

I want to know how many Students I create

```
class Student {
    int assns, mt, final;
    static int count; // static - associated with the class - not one per
                      // object
public:
    Student( ... ): ... { ++ count; }
    // accessors.
    static int getCount() { return count; }
    // static methods - have no 'this' parameter.
    // - not really methods - scope functions
};

int Student::count = 0; // must define the variable (.cc file)
Student s1 {...}, s2 {...}, s3 {...};
cout << Student::getCount(); // 3
```

Now I want to count objects in other classes. How can we abstract this solution into reasonable code?

```
template <typename T> struct Count {
    static int count;
    Count () { ++count; }
    Count (const Count &) { ++count; }
    Count (Count &&) { ++count; }
    ~Count() { --count; }
    static int getCount { return count; }
};

template <typename T> int Count <T>::count = 0;
```

```
class Student: Count<Student> { // private inheritance - inherits Count's
    implementation without creating is-a relationship.
    // - members of Count become private in Student
    int assns, mt, final;
public:
```

```
Student( ... ): ... { ... }  
    // accessors  
    using Count::getCount; // make Count::getCount visible  
};  
  
class Book: Count<Book> { ...  
public:  
    using Count::getCount;  
};
```

Why is `Count` a template?

So that for each class `C`, `class C: Count<C>` creates a new, unique instantiation of `Count`, for each `C`.

This gives `C` its own counter, vs. sharing one counter over all classes.

This technique - inheriting from a template specialized by yourself - looks weird, but happens enough to have its own name - the Curiously Recurring Template Pattern (CRTP).

What is CRTP good for?

Problem 27: Resolving Method Overrides At Compile-Time

Recall: Template Method Pattern

```
class Turtle {
public:
    void draw {
        drawHead();
        drawShell();
        drawTail();
    }
private:
    void drawHead();
    virtual void drawShell()=0; // vtable lookup
    void drawTail();
};

class RedTurtle: public Turtle {
    void drawShell() override;
};
```

Consider

```
template <typename T> class Turtle {
public:
    void draw() {
        drawHead();
        static_cast <T*>(this)->drawShell();
        drawTail();
    }
private:
    void drawHead();
    void drawTail();
};

class RedTurtle: public Turtle<RedTurtle> {
    friend class Turtle;
    void drawShell();
};
```

- no virtual methods, no vtable lookup

```
class GreenTurtle: public Turtle<GreenTurtle> {  
    friend class Turtle;  
    void drawShell();  
};
```

Drawback - no relationship between RedTurtle and GreenTurtle. \implies can't store a mix of them in one container.

Can give Turtle a parent:

```
template <typename T> class Turtle: public Enemy { ... };
```

Then we can store RedTurtles with GreenTurtles

- But can't access draw method
- Could give Enemy a virtual draw.

```
class Enemy {  
public:  
    virtual void draw()=0;  
};
```

- but then there will be a vtable lookup
- On the other hand, if `Turtle::draw` calls several would be virtual helpers, could trade several vtables lookups for one.

Problem 28: Polymorphic Cloning

```
Book *pb = ...;
Book *pb2 = ...; // I want an exact copy of *pb
```

Can't call a ctor directly - don't know what `*pb` is \implies which ctor to call.

Standard solution virtual clone method

```
class Book {
    ...
public:
    virtual Book *clone { return new Book { *this }; }
};

class Text: public Book {
    Text *clone override { return new Text { *this }; }
};

// comic - similar
```

Can we reuse code?

Works better with an abstract base class:

```
class AbstractBook {
public:
    virtual AbstractBook clone()=0;
    virtual ~AbstractBook();
};

template <typename T> class Book_cloneable: public AbstractBook {
public:
    T *clone() override { return new T { static_cast <T*>(*this) }; }
    // T should be AbstractBook to make it compile
};
```

```
class Book: public Book_cloneable<Book> { ... };  
class Text: public Book_cloneable<Text> { ... };  
class Comic: public Book_cloneable<Comic> { ... };
```

If you are curious, see [enable_shared_from_this](#).

Problem 29: Logging

Want to encapsulate “logging” functionality and “add” it to any class.

```
template <typename T, typename Data> class Logger {
public:
    void loggedSet (Data x) {
        cout << "Setting_data_to_" << x << endl;
        static_cast<T *>(this)->set(x);
    }
};

class Box: public Logger <Box, int> {
    friend class Logger;
    int x;
    void set (int y) { x = y; }
public:
    Box(): x{0} {loggedSet(0)};
};

Box b;
b.loggedSet(1);
b.loggedSet(4);
b.loggedSet(7);
```

Another approach:

```
class Box {
    int x;
public:
    Box(): x{0} {}
    void set(int y) { x=y; }
};

template <typename T, typename Data> class Logger: public T {
public:
    void loggedSet(Data x) {
        cout << "Setting_data_to_" << x << endl;
        set(x);
    }
}
```



```
};  
  
using BoxLogger = Logger <Box, int>;  
  
BoxLogger b;  
b.loggedSet(1);  
b.loggedSet(4);  
b.loggedSet(7);
```

- again - no virtual method overhead.

33.1. Mixin inheritance

Note if SpecialBox is a subclass of Box, then Logger <Box, int> is not a subclass of SpecialBox in the second solution. - Also Logger <Box, int> is not a subtype of Logger <Box, int>.

But in the first solution - SpecialBox is a subtype of Logger <Box, int> \implies can specialize behaviour via virtual functions.

Problem 30: Total Control

C++ lets you control how pretty much everything is done - copies, parameter passing, initialization, method call resolution, etc.

Why not go the rest of the way and take control over memory allocation?
Memory allocation are tricky to write. So 2 questions: Why? and How?

Why might you need a custom allocator?

- built-in one is too slow.
 - general purpose - not optimized for any specific use.
 - e.g. if you know you will always allocate objects of the same size, a custom allocator may perform better.
- optimize locality
- to use “special memory”
- to profile your program (collect stats)

How do you customize allocation?

- overload operator `new`
- if you write a global operator `new`, all allocations in the program will use your allocator
- also write operator `delete` - else undefined behaviour.

Eg

```
void *operator new(size_t size) {
    cout << "Request for " << size << "bytes\n";
    return malloc(size);
}
void operator delete(void *p) {
    cout << "Freeing " << p << endl;
    free(p);
}

int main() {
    int *x = new int;
    delete x;
```

```
}
```

Works

- but is not correct. Does not adhere to connection. If operator `new` fails, it is supposed to throw `bad_alloc`.
- actually, if operator `new` fails, it is supposed to call the `new_handler` function.
 - the `new_handler` function can
 - free up space (somehow)
 - install a different `new_handler`/ deinstall the current.
 - throw `bad_alloc`
 - abort/exit
- `new_handler` should be called in an ∞ loop.
- if `new_handler` is `nullptr`, then operator `new` throw.
- Also - `new` must return a valid ptr if `size==0`, and `delete(nullptr)` must be safe.

Corrected operator `new`:

```
#include <new>
void *operator new(size_t size) {
    cout << "Request for " << size << "bytes\n";
    while (true) {
        void *p = malloc(size);
        if (p) return p;
        std::new_handler h = std::get_new_handler();
        if (h) h();
        else throw std::bad_alloc {};
    }
}

void operator delete(void *p) {
    if (!p) return;
    cout << "Freeing " << p << endl;
    free(p);
}
```

Replacing global operator `new/delete` affects your entire program.
Probably want to replace these operators on a class-by-class basis.
- especially if optimizing for the size of your objects.

Define operator `new/delete` within the class - must be static members.

```
class C {
public: ...
    static void *operator new(size_t size) {
        cout << "Running C's allocator\n";
        return ::operator new(size);
    }
}
```

```
static void operator delete(void *p) noexcept {
    cout << "Freeing_" << p << endl;
    return ::operator delete(p);
}
};
```

Can we write to an arbitrary stream?

```
class C {
public:
    static void *operator new(size_t size, std::ostream &out) {
        out << ... << endl;
        return ::operator new(size);
    }
    ...
};

C *x = new(cout) C;
ofstream f {...};
C *y = new(f) C;
```

Must also write the corresponding delete:

```
class C {
...
    static void operator delete(void *p, std::ostream &out) noexcept {
        out << ... << endl;
        return ::operator delete(p);
    }
};
```

Won't compile! You also need ordinary delete.

```
class C {
...
    static void operator delete(void *p, std::ostream &out) noexcept {
        out << ... << endl;
        return ::operator delete(p);
    }
};

C *p = new (f) C; // ofstream allocator
delete p; // ordinary delete
```

- Must have ordinary delete - else compile error.

How can you call the specialize one? You can't.

Then why do you need it?

If the ctor that runs after specialized new - throws - then specialized operator delete is called.

If there isn't one \Rightarrow no delete is called \Rightarrow leak

Problem 31: Total Control ove Vectors & Lists

Incorporating custom allocators into containers.

Make the allocator an argument to the template. - With a default value.

```
template <typename T, typename Alloc = Allocator<T>> class Vector { ... };

template <typename T> class allocator {
public:
    // big 5 (no fields)
    T *address(T &x) { return &x; }
    T *allocate(size_t n) { return ::operator new(n*sizeof(T)); }
    void deallocate(T *p) { ::operator delete(p); }
    void construct, destroy // placement versions
}
```

Lists - not so easy

- same Alloc parameter
- want to allocate Node, not T's

```
template <typename T, typename Alloc = Allocator<T>> class lists {
    typename Alloc::rebind<Node>::other alloc;
    // same as before
};

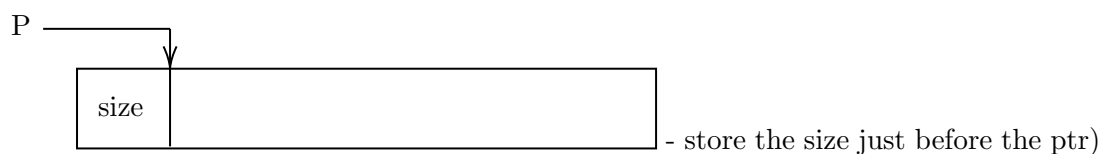
template <typename T> class allocator {
public:
    ...
    template <typename U> struct rebind {
        using other = allocator<U>;
    };
};
```

Problem 32: A fixed-size allocator

Let's write one.

Expect: faster than built-in.

- no need to keep track of sizes (aside: many traditional allocators:



- saves space (no hidden size field)
- saves time (no hunting for a block of the right size)

Create a pool of memory - an array large enough to hold n T objects.

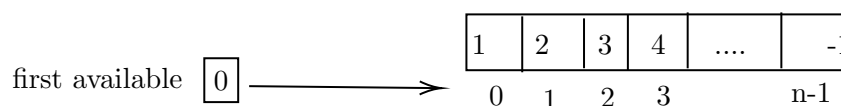
			...	
--	--	--	-----	--

When a slot in the array is given to the client, it will act as a T object.

When we have it, it will act as a node in a linked list.

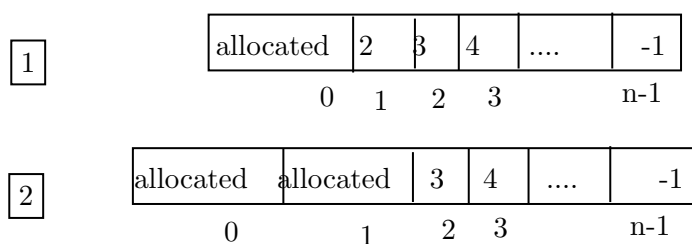
Store an int in each slot - index (not ptr) of the next free slot. Store the index of the first slot.

Initially

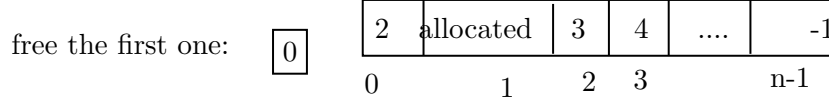


Each node store the index of the next one.

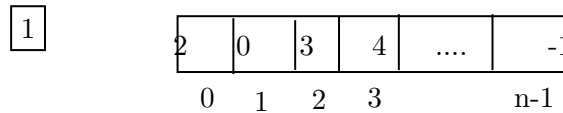
Allocation from the front



Deallocation: add to front of list



free the second one:



Alloc/dealloc: both constant time & very fast

Implementation:

```
template <typename T, int n> class FixedSizeAllocator {
    union Slot {
        int next;
        T data;
        Slot(): next{0} {}
    };
    // Slot - large enough to hold a T. - but also usable as an int.
    Slot theSlots[n];
    int firstAvailable = 0;
public:
    FixedSizeAllocator() {
        for (int i = 0; i < n-1; ++i) theSlots[i].next = i + 1;
        theSlots[n-1].next = -1;
    }
    T *allocate() noexcept {
        if (firstAvailable == -1) return nullptr;
        T *result = &(theSlots[firstAvailable].data);
        firstAvailable = theSlots[firstAvailable].next;
        return result;
    }
    void deallocate(void *item) noexcept {
        int index = (static_cast<char*>(item) - reinterpret_cast<char*>(
            theSlots))/sizeof(Slot);
        theSlots[index].next = firstAvailable;
        firstAvailable = index;
    }
};
```

```
class Student final {
    int assns, mt, final;
    static FixedSizeAllocator<Student, SIZE/* How many slots do you want? */>
        pool;
public:
```

```

Student( ... ) : ...
static void *operator new(size_t size) {
    if(size != sizeof(Student)) throw std::bad_alloc{};
    while (true) {
        void *p = pool.alloc();
        if (p) return p;
        auto h = std::get_new_handler();
        if (h) h();
        else throw std::bad_alloc{};
    }
}
static void operator delete (void *p) noexcept {
    if (!p) return;
    pool.deallocate(p);
}
};
fixedSizeAllocator<Student, SIZE> pool;

// main

int main() {
    Student *s1 = new Student;
    Student *s2 = new Student;
    // custom allocator
    delete s1;
    delete s2;
    // custom deallocator
}

```

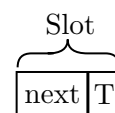
Q where does the memory for s1, s2 reside?

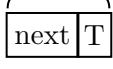
A static Memory (NOT the heap)

- Depending on how you build your allocator, the pool can be in static memory, on the stack, or on the heap.

Notes

- use a union to treat a slot as both an `int` and a T object
- ensures no memory is wasted on bookkeeping.



- alternative: store “next” index adjacent to the T object, i.e.  i.e. struct instead of union.
- advantage of a union: no wasted memory.
- disadvantage: if you access a dangling T ptr, you can corrupt the linked list.

eg

```
Student *s = new Student;
delete s;
s->setAssns( ... ); // will probably corrupt the list
```

(Lesson: following dangling ptrs can be VERY dangerous)

With a strut, 'next' filed is before the ptr, so you have to work harder to corrupt it: e.g. `reinterpret_cast<int*>(s)[-1]=...`

Unions - if one field has a ctor, the union needs a ctor - since it's a union, the ctor should only initialize one filed.

On the other hand - if you have a struct - you will have difficulty if T has no default ctor.

Eg

```
struct Slot {
    int next;
    T data;
};
Slot theSlots[n]; // can't do this if T has no default ctor
```

Can't do operator new/placement new - we're writing operator new!

Also - Why indexes vs ptrs?

- `ints` are smaller that ptrs
- So we waste no memory as long as T is at least as large as an `int`.
- We do waste memory if T is smaller that an `int`.
- Could use short, char for the index (if this allows enough items)
- Could make the index type a parameter of the allocator template.

Why is class Student final?

- fixed-size allocator
- subclass may be larger - won't wok with our allocator, but would still use it.
- options
 1. make the class final
 2. check the size, throw if the size is not right
 3. check the size, use global operator new/delete if size is not right
 4. give subclass its own allocator

Problem 34: I want a (tiny bit) smaller vector class

`vector/vector_base` have an allocator field which itself has no fields.

What is its size? 0? No - not allowed in C++. \therefore all types size > 1 .

Practically - compiler inserts a dummy character.

So the allocator make the vector one byte larger. (possibly more, due to alignment)

To save this space - C++ provides: empty base optimization (EBO).

- an empty base class does not have to occupy space in an object.

So make the allocator a base class of vector. At the same time, make `vector_base` a super-class.

```
template <typename T, typename Alloc=allocator<T>> struct vector_base: Alloc
{
    size_t n, cap;
    T *v;
    using Alloc::allocate;
    using Alloc::deallocate;
    // etc
    // Because we don't know anything about Alloc, we need to bring these
    // members into scope.
    ...
};
template <typename T, typename Alloc=allocator<T>> class vector: private /*
    no is-a*/ vector_base <T, alloc> {
    using vector_base<T, alloc>::n;
    using vector_base<T, alloc>::cap;
    using vector_base<T, alloc>::v;
    ...
};
```

Details are left as exercise.

37.1. Something Missing from Last Year

So we can eliminate the space cost of an allocator by making it a base class. At the same time, make `vector_base` a base class of a `vector`.

```

template <typename T, typename Alloc = allocator<T>>
struct vector_base: private Alloc { // struct has default public inheritance
    size_t n, cap;
    T *v;
    using Alloc::allocate;
    using Alloc::deallocate;
    // etc.
    vector_base(size_t n): n{0}, cap{n}, v{allocate(n)} {}
    ~vector_base() { deallocate(v); }
};

template <typename T, typename Alloc = allocator<T>>
class vector: vector_base<T, Alloc> { // private inheritance - no is-a
    relation
    using vector_base<T, Alloc>::n;
    using vector_base<T, Alloc>::cap;
    using vector_base<T, Alloc>::v;

    using Alloc::allocate; // or say this->allocate
    using Alloc::deallocate; // this->deallocate
public:
    ... Use n , cap, v instead of vb.n, vb.cap, vb.v
};

```

uninitialized_copy, etc. - need to call construct/destroy
 - Simplest - let the take an allocator as a parameter

```

template <typename T, typename Alloc> {
    void uninitialized_fill(T *start, T *finish, const T &x, Alloc a) {
        ...
        a.construct(...)
        ...
        a.destroy(...)
        ...
    }
}

```

How can vector pass an allocator to these functions?

```

uninitialized_fill(v, v + n, x, static_cast<Alloc&>(*this)); // Cast
yourself to base class reference

```

Remaining details - exercise

Part IV.

Other

38. Tutorials	149
38.1. Bash	149
38.2. Debugger	152
38.2.1. GDB Commands	152
38.3. Git	152
38.3.1. What is version control?	152
38.3.2. Using git: the basics	153
38.3.3. Working on multiple computers: setting up git	153
38.3.4. Preventing committing the wrong files	154
38.3.5. Merge conflicts	154
38.3.6. Retrieve old versions of files	154
38.3.7. Collaborating and rolling back more effectively: branches	155
38.3.8. ...and much more	155
38.4. Parsing	155
38.5. MVC	156
38.5.1. Ncurses	156
39. Readings	157

38.1. Bash

```
$ ls -list
file1 file2
$ pwd
/u/snwhitch/cs246e/1189/a1
$ cat file1 file2
This is file1
This is file2
$ echo file1 file2
file1 file2
$ rm file2
$ ls
file2
$ cp file2 file1
$ ls
file1 file2
$ mv file2 file3
$ ls
file1 file3
$ touch file2
$ ls
file1 file2 file3
$ touch file3 -- change timestamp
$ mktemp
/tmp/tmp ABCD
$ mkdir dir
$ cd dir
$ diff file1 file2
$ echo $?
$ wc file1
x y z file1
$ less file1
$ man ls
$ head -10 /usr/share/dict/words
$ tail -10 /usr/share/dict/words
$ tr ' ' ':'
```

```
$ ls
file1 file2
$ ls | tr ' ' '\n'
file1
file2
$ ls | tr ' ' '\n' | wc -l
2
$ wc < file
$ cat args
-la
$ ls $(cat args)
equal to
$ ls -la
```

In today's tutorial we solved the following problem: Given a filename as input, create copies of that file until exactly 10 in the current directory have that content.

First, we used `system()` to create copies `copy1`, ..., `copy9`, regardless of what files were in the directory. The code was as follows:

```
#include <iostream>
#include <string>

using namespace std;

// Run command s, return the exit code
int sys(string s) {
    int res = system(s.c_str());
    return WEXITSTATUS(res);
}

int main(int argc, char *argv[]) {
    string filename = argv[1];

    for (char c = '1'; c <= '9'; c++) {
        sys("cp" + filename + "_" + c);
    }
}
```

We extended this to a second program that solved the original problem:

```
#include <iostream>
#include <string>

using namespace std;
```

```
// Run command s, return the exit code
int sys(string s) {
    int res = system(s.c_str());
    return WEXITSTATUS(res);
}

// Check if the contents of two files are the same
bool samefile(string file1, string file2) {
    return sys("diff " + file1 + " " + file2) == 0;
}

int main(int argc, char *argv[]) {
    string filename = argv[1];

    int copies = 0; // number of existing copies
    FILE *f = popen("ls", "r");
    char buf[128];
    while (fscanf(f, "%s", buf) != EOF) {
        string newfile = buf;
        if (samefile(filename, newfile)) copies++;
    }
    for (char c = '1'; c <= '0' + 10 - copies; c++) { // e.g. when copies =
        sys("cp " + filename + " " + c);
    }
}
```

There was also a good question asked in the tutorial about reading/writing with pipes. To answer the first question: pipes do not have to be bidirectional on any given system. You can get something to that effect with two pipes in alternate directions, however.

Additionally, you can use the file returned by `popen` in write mode to provide input to the command in question. For example,

```
#include <stdio>
#include <stdlib>

int main() {
    FILE *f = popen("wc", "w");
    fprintf(f, "Test");
    pclose(f);
}
```

and you will get output

```
0 1 4
```

as expected (note `wc` counts newline characters to determine a line count, and we gave it none)

38.2. Debugger

```
$ valgrind ./a.out
```

definitely lost

-g no optimization and compile with extra info

38.2.1. GDB Commands

q	quit	
r	run	
b	break	set a break point
n	next	go to next line
p	print	print an expression
b+	backtrace	print the call stack
up/down		move up/down call stack
l	list	print code
watch		set a watch point
c	continue	resume running

38.3. Git

```
$ git init
$ git add Hello
$ git status
$ git commit -m "sth"
$ git commit
$ git rm Hello
$ git log
$ git push origin master
$ git clone
$ git log
$ git check out 7914e88e98a898a889 test.cc
$ git reflog
$ git reset HEAD@{3}
$ git checkout newbranch
```

38.3.1. What is version control?

Often when writing code we want to be able to look back at old versions of code, whether because we made a mistake when changing something and want to roll back, or simply want to remember what code used to look like. Similarly, we want the ability to manage multiple versions of a codebase (owned by different people), and to synchronize them between computers. Version control helps with these tasks. There have been a number of popular version control tools over the years, including (but not limited to)

- Subversion (svn)

- Git
- Mercurial
- Concurrent Versions System (cvs)

Currently git is arguably the most famous and popular, in no small part due to the support of github.com, but all have seen plenty of use. We will focus on git for this tutorial.

38.3.2. Using git: the basics

To use git you'll need a repository. While we could use the course repository (1189), you don't have permissions to do some operations on it, so we'll make our own.

1. First, navigate into whatever folder you want your repository to live in. Type `git init` to create an empty git repository.
2. At this point, create some files. Once you've done so, you can ask git to start tracking them (or notice an update of a file that's changed) using `git add`. If you `git add` a folder, it will recursively add every file in the folder.
3. Once you are ready to store a revision (version) of your code long-term, use `git commit` to have git remember the current state of any files you added to it since your last commit. You should supply a short "commit message" so that someone looking at your commits later can tell what changed here. To do this, you can use `git commit -m 'Your message here'`. If you don't use `-m`, git will open `$EDITOR` and ask you to enter a commit message.
4. If you want to move a file to a different folder and have git recognize this, delete a file from the repository, or any other usual filesystem task, use "`git mv`", "`git rm`", etc. Note that these both perform the actual action you'd expect (moving the file, deleting the file, etc) while also telling git that it's happened so that the repository can be updated next time you commit.

38.3.3. Working on multiple computers: setting up git

You can share a repository among multiple computers as long as one of them hosts the repository: it's easiest to do this by using a website such as github. For assignments it's best not to use github unless you have access to private repositories since otherwise you may inadvertently share your assignment solutions with other students, but luckily UW has their own service (UW gitlab) at <https://git.uwaterloo.ca/> for this purpose. You can have up to 10 private repositories at a time, which should be more than enough.

1. First, try creating a repository on gitlab.
2. You can add a remote repository to your machine by using `git clone` followed by the repository's URL. Most services (such as github and gitlab) should give you a link to copy. Try cloning your repository on to your machine.
3. You can send your commits to any other people using your repository by using `git push`. Try creating some files, committing them, and then pushing them. You should now be able to see them on gitlab.
4. If you want to receive any changes other users have made, you can use `git fetch` or `git pull` to receive other peoples' changes. We'll talk about the difference later. For now, try cloning the repository into a second location and using `git push/pull` to share some files.

38.3.4. Preventing committing the wrong files

There are lots of files you may not want to track in your git repository. For example, usually you don't want to track pdfs and intermediate files related to tex documents, binaries, object files, dependency files, temporary vim/emacs files and so forth. While we could just avoid adding them using 'git add', unfortunately these files will still be added if we add directories. Luckily there's a more convenient way to avoid adding certain files you don't want. Create a file called '.gitignore' in the root directory of the project and commit it. It should contain a list of files to ignore, one per line. You can use globbing patterns, so for example you can tell git to ignore all object files by putting '*.o' in the .gitignore file. Try creating a .gitignore file which excludes vim swap files (files beginning with a . and ending in .swp) and the file "a2q3b", then test it out by putting some files in a directory, some of which are swap files or named a2q3b, and adding them to the repository.

38.3.5. Merge conflicts

Once you have two or more people working on the same project you may start running into merge conflicts. Even if you're the only person working on your project, if you are working on multiple machines you may occasionally forget to commit or push changes, which may result in merge conflicts anyways. Put simply, a merge conflict occurs when the code in your repository disagrees with the repository you're pulling from.

1. Try changing a file on two copies of your repository. Commit and push one of them. On the other, don't add or commit it, and try using git pull.
2. Commit your file and try using git pull again. Git will merge the two versions and commit histories, which will likely require you to manually synchronize any files which were modified by each side.
3. Alternatively, you can use git fetch to receive remote changes without alternatively attempting to merge. If you do, you will have to manually merge with git merge afterwards. You must merge before you can git push your commits to other copies of the repository.

Merging can be annoying and occasionally difficult. Whenever possible you should try to commit early and often, pull before you make changes, and try not to work on the same file between multiple machines or people unless absolutely necessary. That being said, sometimes it's unavoidable, and accidents happen...

38.3.6. Retrieve old versions of files

Occasionally you may want to roll back parts of your code (or all of it) to a previous version, whether because you created a bug which you can't identify, because you deleted a file and now need it back, or for any number of other reasons.

1. Delete a file from your repository using git rm. Commit and then make another few changes and commit each time.
2. Now we want to get the file back (the same process works to restore old versions of existing files). View the recent commit history using git log.
3. Once you've found the commit you're looking for (this is one of the places where commit messages come in handy!), note its commit number. We can now check out the file we want from the version we want via git checkout [commit number] [file]. This doesn't change the state of the repository: you still need to re-add and commit it if you want the changes to be recognized long-term.

38.3.7. Collaborating and rolling back more effectively: branches

As you can see, merging and retrieving old files both work, but they're a bit clunky. Constant merging when collaborating with someone (or alternatively coordinating to make sure you're never changing the same files) gets annoying and time-consuming and can be a major source of bugs. Identifying which commit holds the version of a file you want can be difficult, especially if you need to roll back a major feature, code refactor or other change. Git provides branches to make this process easier.

You may have noticed that everything git has said so far tends to have the word "master" on it somewhere or other. This is the default branch. A branch is essentially a version of the repository: anyone on the same repository has access to every branch, but at any one time you're only working on a single branch and your changes don't affect anyone working on a different branch than you are.

1. Create a new branch of your repository using `git branch [branch name] [old branch name]`, which creates a new branch starting from the old branch. You can move to the branch using `git checkout [branch name]`.
2. Try making some changes to the new branch, committing them and pushing them.
3. In the other copy of your repository, make some changes to master, commit them and push them. Notice how you have no merge conflicts.
4. Back on the branch, pull all the changes made to master, then use `git merge [branch name]` to merge your branch with master. This lets you do all of the merging at once and have well-defined places where features are incorporated into the main codebase. In serious software projects all development is strongly encouraged to happen on branches rather than on master: whether you choose to follow this guideline is up to you.

38.3.8. ...and much more

Git has many more commands, features, options to commands and so forth than we have time to explore in this short tutorial. You can find plenty of tutorials online, and can view git's manual using `man git` and `man git-[command]` (for example, `man git-checkout`). You should try using git or another VCS to manage your code in CS246E from now on as well as your personal projects, since version control is going to be a mainstay of your life from here on out!

38.4. Parsing

```
// bool2.h
#include <string>

using namespace std;

class expression {
public:
    virtual pretty_print() = 0;
    virtual bool eval() = 0;
    virtual void set (string name, bool val) = 0;
};

class or_expr: public expression {
```

```

    unique_ptr<expression> left, right;
public:
    or_expr(...): left{move(left)}, right(...) {}
    string pretty_print() override {
        return ...
    }
    ...
};

class and_expr: public expression {...};

class variable: public expression {
    string name; bool val;
public:
    variable(string name, bool val=false): name(name), val(val) {}
    ...
};

unique_ptr<expression> parse_and(const string &s);

unique_ptr<expression> parse_or(const string &s) {
    size_t i = s.find('|');
    if (i == string::npos) return parse_and(s);
    string s1 = s.substr(0, i);
    string s2 = s.substr(i+1);

    return make_unique<or_expr> (parse_and(s_1), parse_or(s_2));
}

```

38.5. MVC

38.5.1. Ncurses

Here are readings relevant to what we've done so far. Note that I am not going to test you on the textbook. But if you need more detail on lecture topics, this would be a place to start.

7.7.1--7.7.4
12.2.5--12.3.5
12.6.1
17.1--17.2.2, 17.2.4, 17.3--17.3.3, 17.4.1, 17.4.4, 17.5.1, 17.5.2, 17.6
16.2.3-4, 16.2.9, 16.2.13, 9.5.1, 30.3.2, 19.4
3.4.1-2, 23.2--23.4.6.1 (omit 23.4.4 for now), 23.5--23.5.3.1
17.3.4
11.2.4 (for placement new; presentation is mixed in with more advanced stuff
; could skip)
28.6.3
13.2
5.2, 34.3.1, 17.6.4, 19.2.3
13.3
13.6
35.5.1, 35.5.2
13.5.1.1-2
31.3.7, 31.3.3
3.4.3, 19.2.2
11.4